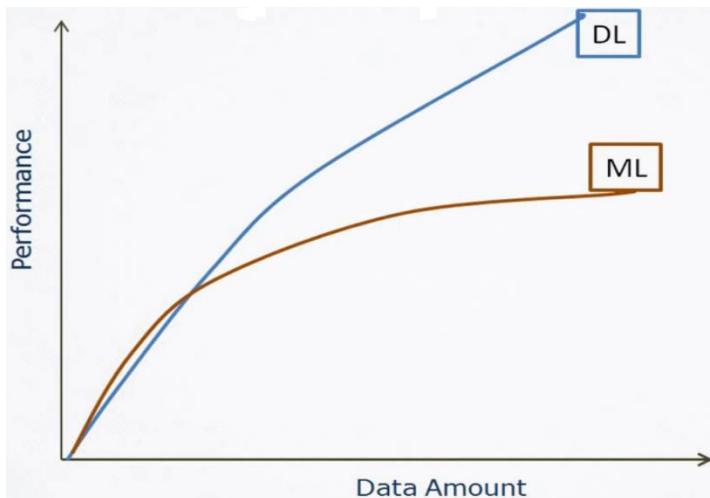
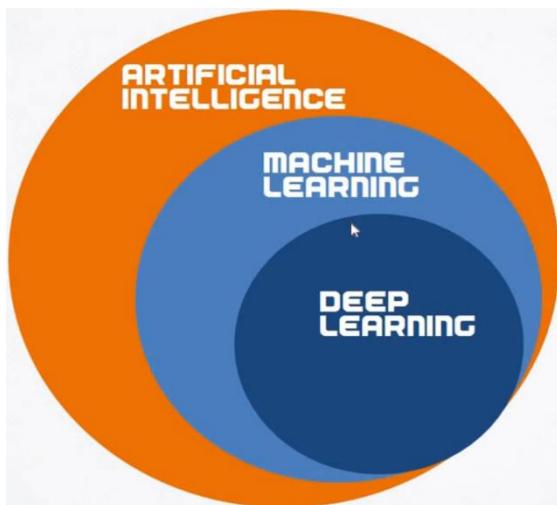
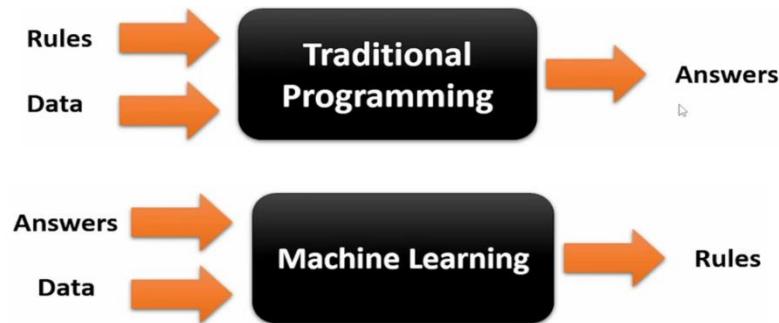


## Deep Learning



Deep learning is a computer software that **mimics the network of neurons in a brain**. It is a subset of machine learning and is called deep learning because it makes use of **deep neural networks**.

## Deep Learning Model Process



### Deep learning use cases



*Self Driving Car*



*Robotics & Labour Automation*



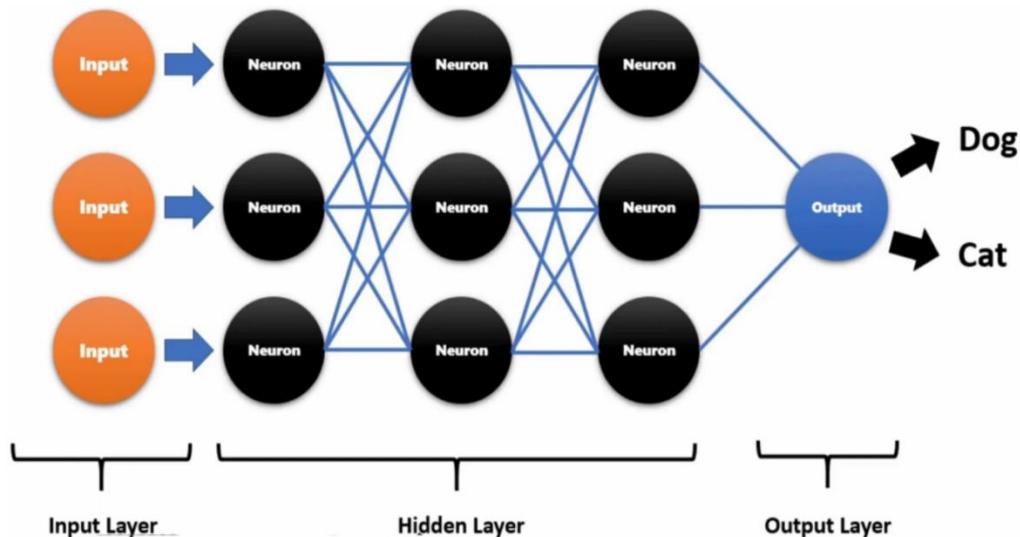
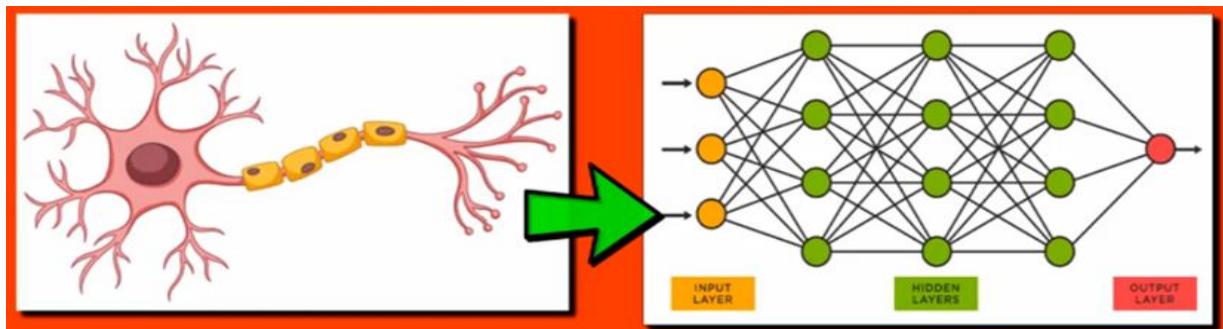
*Tumour Detection*

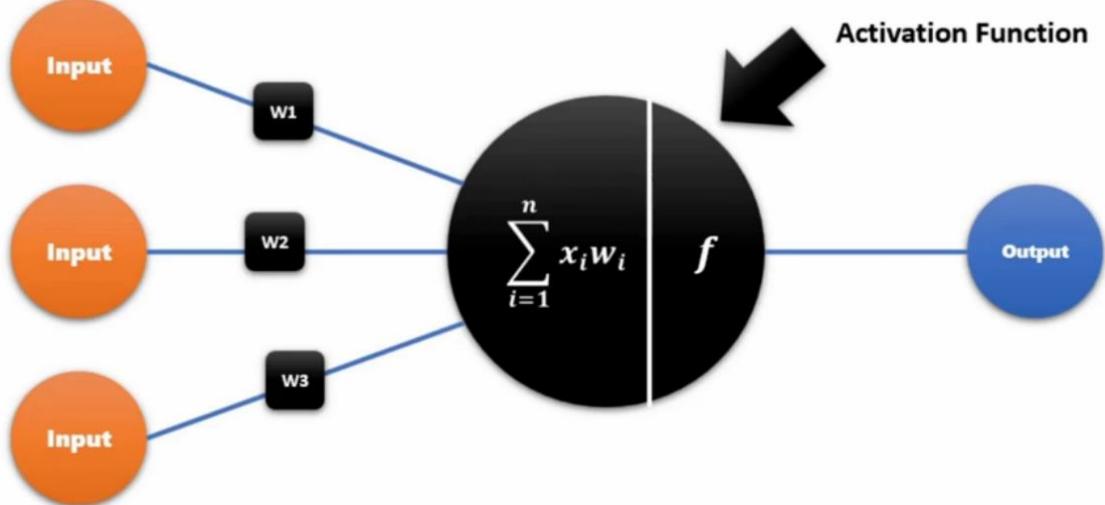
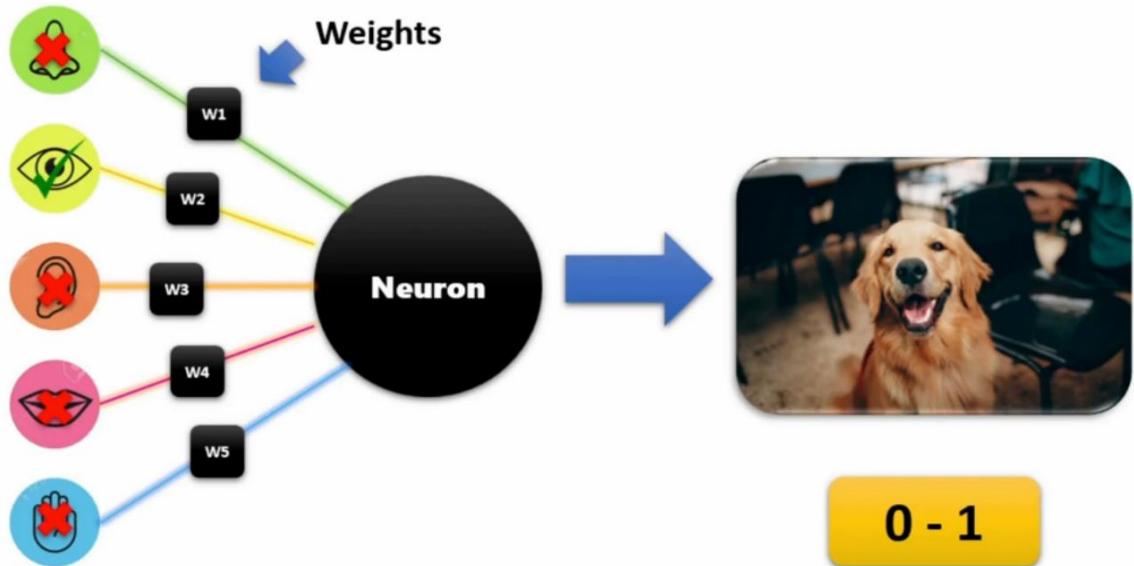


*Speech Recognition*

Deep learning algorithms are constructed with connected layers.

- ✓ The first layer - Input Layer
- ✓ The last layer - Output Layer
- ✓ All layers in between are called Hidden Layers.





**Formula** → Activation Function( $\sum_{i=1}^n x_i w_i + b$ )

**Types** → Linear and Non - Linear

1. Load data
2. define keras models
3. compile
4. fit / train
5. evaluate

```
#(1)Load data
import numpy as np

dataset = np.loadtxt('/content/pima-indians-diabetes.data.csv', delimiter=',')
X = dataset[:,0:8]
y = dataset[:,8]
```

`dataset`

```
array([[ 6.   , 148.   , 72.   , ...,  0.627,  50.   ,  1.   ],
       [ 1.   , 85.   , 66.   , ...,  0.351,  31.   ,  0.   ],
       [ 8.   , 183.   , 64.   , ...,  0.672,  32.   ,  1.   ],
       ...,
       [ 5.   , 121.   , 72.   , ...,  0.245,  30.   ,  0.   ],
       [ 1.   , 126.   , 60.   , ...,  0.349,  47.   ,  1.   ],
       [ 1.   , 93.   , 70.   , ...,  0.315,  23.   ,  0.   ]])
```

`X`

```
array([[ 6.   , 148.   , 72.   , ..., 33.6   ,  0.627,  50.   ],
       [ 1.   , 85.   , 66.   , ..., 26.6   ,  0.351,  31.   ],
       [ 8.   , 183.   , 64.   , ..., 23.3   ,  0.672,  32.   ],
       ...,
       [ 5.   , 121.   , 72.   , ..., 26.2   ,  0.245,  30.   ],
       [ 1.   , 126.   , 60.   , ..., 30.1   ,  0.349,  47.   ],
       [ 1.   , 93.   , 70.   , ..., 30.4   ,  0.315,  23.   ]])
```

`y`

```
array([1., 0., 1., 0., 1., 0., 1., 0., 1., 0., 1., 1., 1., 1.,
       1., 0., 1., 0., 0., 1., 1., 1., 1., 0., 0., 0., 0., 1., 0., 0.,
       0., 0., 0., 1., 1., 0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 0.,
       1., 1., 0., 0., 0., 1., 1., 0., 1., 0., 1., 0., 0., 0., 1.,
       ...,
       0., 0., 0., 0., 1., 1., 0., 0., 1., 0., 1., 1., 0., 0., 1., 0.,
       1., 1., 0., 0., 1., 1., 0., 1., 0., 1., 0., 1., 0., 0., 0.,
       0., 1., 0.])
```

*Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...*

```
#(2)Define keras models
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(12,input_dim=8,activation='relu'))
model.add(Dense(8,activation='relu'))
model.add(Dense(1,activation='sigmoid'))
```

```
#(3)Compile
model.compile(loss='binary_crossentropy')
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 12)	108
dense_1 (Dense)	(None, 8)	104
dense_2 (Dense)	(None, 1)	9

Total params: 221  
Trainable params: 221  
Non-trainable params: 0

```
#(4)Fit/Train
model.fit(X,y,epochs=150,batch_size=10)
```

Epoch 1/150  
77/77 [=====] - 1s 2ms/step - loss: 4.5729  
Epoch 2/150  
77/77 [=====] - 0s 2ms/step - loss: 1.1170  
77/77 [=====] - 0s 2ms/step - loss: 0.7077  
Epoch 13/150  
...  
Epoch 149/150  
77/77 [=====] - 0s 2ms/step - loss: 0.5202  
Epoch 150/150  
77/77 [=====] - 0s 2ms/step - loss: 0.5245  
Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output [settings](#)...

```
#(5)Evaluate
model.evaluate(X,y)
```

24/24 [=====] - 0s 1ms/step - loss: 0.4989

0.49893298745155334

```
[2] import tensorflow as tf  
import numpy as np  
from tensorflow import keras
```

```
[3] print(tf.__version__)
```

2.6.0

## Data Set

```
▶ xs = np.array([-1.0, 0.0, 1.0, 2.0, 3.0, 4.0], dtype=float)  
ys = np.array([-3.0, -1.0, 1.0, 3.0, 5.0, 7.0], dtype=float)
```

## Simple Neural Network

```
▶ model = tf.keras.Sequential([keras.layers.Dense(units=1, input_shape=[1]))
```

## Compiling

```
[6] model.compile(optimizer='sgd', loss='mean_squared_error')
```

## Training the Neural Network

```
▶ model.fit(xs, ys, epochs=500)  
  
▶ Epoch 495/500  
1/1 [=====] - 0s 7ms/step - loss: 3.1498e-05  
Epoch 496/500  
1/1 [=====] - 0s 6ms/step - loss: 3.0851e-05  
Epoch 497/500  
1/1 [=====] - 0s 12ms/step - loss: 3.0217e-05  
Epoch 498/500  
1/1 [=====] - 0s 8ms/step - loss: 2.9597e-05  
Epoch 499/500  
1/1 [=====] - 0s 5ms/step - loss: 2.8990e-05  
Epoch 500/500  
1/1 [=====] - 0s 5ms/step - loss: 2.8394e-05  
<keras.callbacks.History at 0x7fda501f5890>
```

## Prediction

```
[8] print(model.predict([5.0]))
```

[[8.995719]]

## Dot Product

Matrix A -> m × n matrix

Matrix B -> n × p matrix

Dot product of matrix A and matrix B = ?

(m × n).(n × p) = m × p matrix

where the n entries across a row of A are multiplied with the n entries down a column of B and summed to generate an entry of resulting matrix.

## Why we do Dot Product

The dot product greatly reduces the computation time especially when we have a large number of independent equations to solve.

Note : number of columns in the first matrix should always be equal to the number of rows in the second matrix.

$$\begin{bmatrix} p_1 & q_1 & r_1 \\ p_2 & q_2 & r_2 \end{bmatrix} \quad \begin{bmatrix} X_1 \\ X_2 \\ X_3 \end{bmatrix}$$

## Dot Product Illustration

$$\begin{bmatrix} p_1 & q_1 & r_1 \\ p_2 & q_2 & r_2 \end{bmatrix} \bullet \begin{bmatrix} X_1 \\ X_2 \\ X_3 \end{bmatrix} = \begin{bmatrix} p_1xX_1 + q_1xX_2 + r_1xX_3 \\ p_2xX_1 + q_2xX_2 + r_2xX_3 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix}$$

Shape = (2×3).(3×1) = (2×1)

## Broadcasting

Matrix A -> m × n matrix

Matrix B -> shape should be shape (1 x n) or (m X 1) or just a scalar number.

Row vector [x<sub>1</sub> x<sub>2</sub> x<sub>3</sub> .. X<sub>n</sub>]

Column vector

$$\begin{bmatrix} x_1 \\ x_2 \\ \cdot \\ \cdot \\ x_n \end{bmatrix}$$

➤ If B shape is (1 x n) matrix (row vector),  
then it gets replicated itself column-wise to become (m x n) matrix.

➤ If B shape is (m x 1) matrix (column vector),  
it gets replicated row-wise to become (m x n) matrix.

➤ If it is a scalar number, then  
it gets converted to (m x n) matrix where each of the element is equal to the scalar number

➤ Broadcasting also works in case of function

### Broadcasting Illustration

$$\begin{bmatrix} p_1 & q_1 & r_1 \\ p_2 & q_2 & r_2 \end{bmatrix} + [x_1 \ x_2 \ x_3] = \begin{bmatrix} p_1 + x_1 & q_1+x_2 & r_1 + x_3 \\ p_2 + x_1 & q_2+x_2 & r_2 + x_3 \end{bmatrix} = \begin{bmatrix} y_1 & y_2 & y_3 \\ y_4 & y_5 & y_6 \end{bmatrix}$$

### 2. Multiplication of a matrix and a column vector

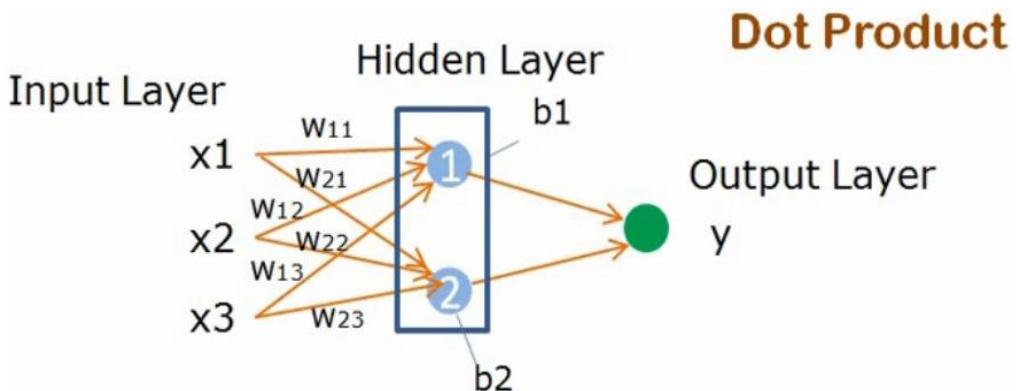
$$\begin{bmatrix} p_1 & q_1 & r_1 \\ p_2 & q_2 & r_2 \end{bmatrix} \times [x_1 \ x_2 \ x_3] = \begin{bmatrix} p_1 \times x_1 & q_1 \times x_2 & r_1 \times x_3 \\ p_2 \times x_1 & q_2 \times x_2 & r_2 \times x_3 \end{bmatrix} = \begin{bmatrix} y_7 & y_8 & y_9 \\ y_{10} & y_{11} & y_{12} \end{bmatrix}$$

### 2. Addition of a matrix and a scalar

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} + 2 = \begin{bmatrix} 1 + 2 & 2 + 2 & 3 + 2 \\ 4 + 2 & 5 + 2 & 6 + 2 \end{bmatrix} = \begin{bmatrix} 3 & 4 & 5 \\ 6 & 7 & 8 \end{bmatrix}$$

### 3. Element-wise function call

$$\begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix} \text{Fun}^{**2} = \begin{bmatrix} 1^{**2} & 2^{**2} & 3^{**2} \\ 4^{**2} & 5^{**2} & 6^{**2} \end{bmatrix} = \begin{bmatrix} 1 & 4 & 9 \\ 16 & 25 & 36 \end{bmatrix}$$



$$Z = W^T X + B$$

$$\hat{Y} = Z$$

$$a = \sigma(Z) = \frac{1}{1+e^{-z}}$$

$$\begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \end{bmatrix} = \begin{bmatrix} w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b_1 \\ w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b_2 \end{bmatrix}$$

$$\begin{bmatrix} w_{11} & w_{12} & w_{13} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + b_1 = w_{11}x_1 + w_{12}x_2 + w_{13}x_3 + b_1 = z_1$$

$$a_1 = \hat{y}_1 = \sigma(z_1)$$

$$\begin{bmatrix} w_{21} & w_{22} & w_{23} \end{bmatrix} \cdot \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} + b_2 = w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + b_2 = z_2$$

$$a_2 = \hat{y}_2 = \sigma(z_2)$$

$$\begin{bmatrix} w_{01} & w_{02} \end{bmatrix} \cdot \begin{bmatrix} a_1 \\ a_2 \end{bmatrix} + b_0 = w_{01}a_1 + w_{02}a_2 + b_0 = z$$

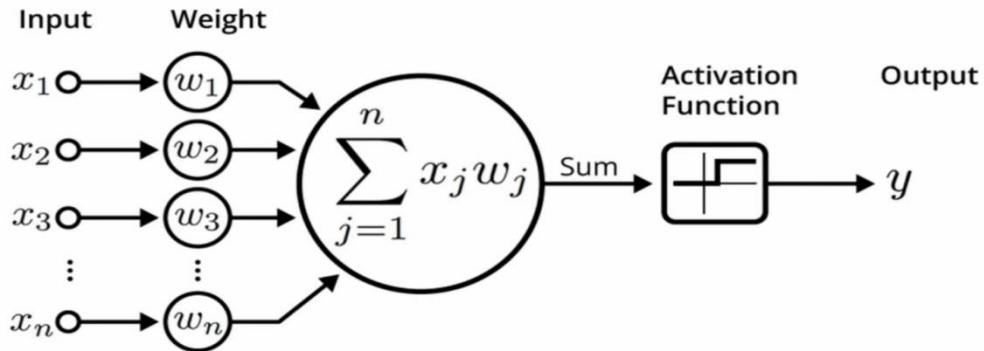
$$y = \sigma(z)$$

## What is Activation Functions?

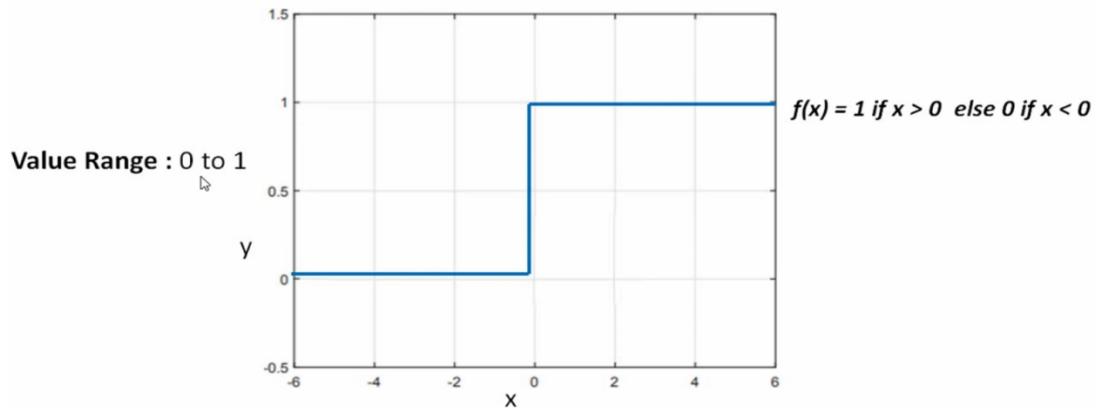
Activation functions are decision making units of neural networks.  
They calculates net output of a neural node

## Why do we need Activation Functions?

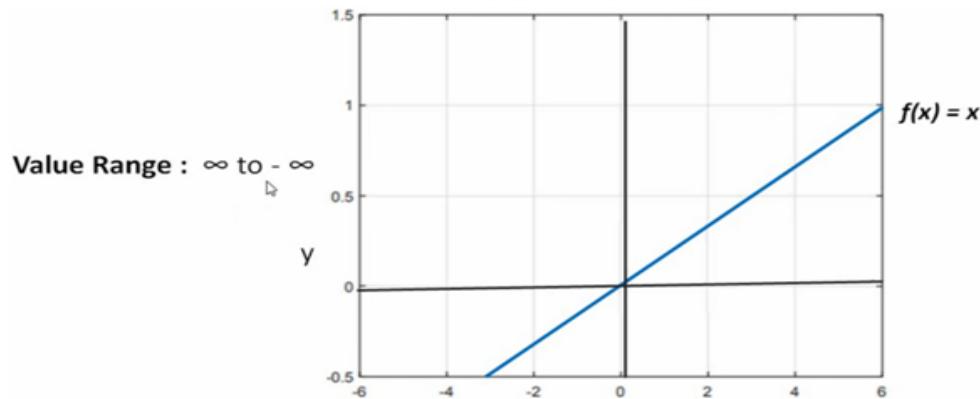
Without activation function neural network is just a linear regression model



### Heaviside or Binary Step Function

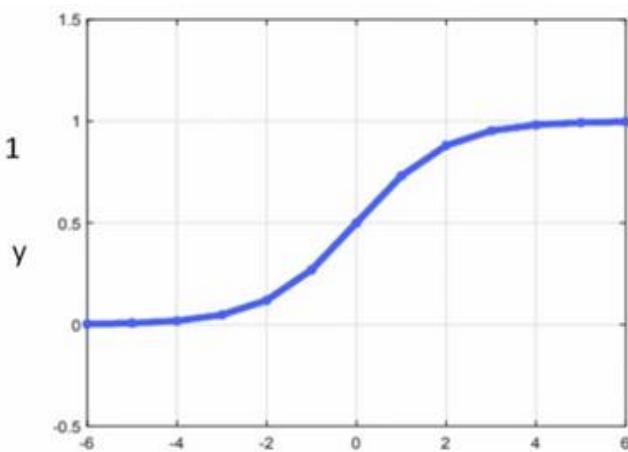


### Linear Function



### Sigmoid Function

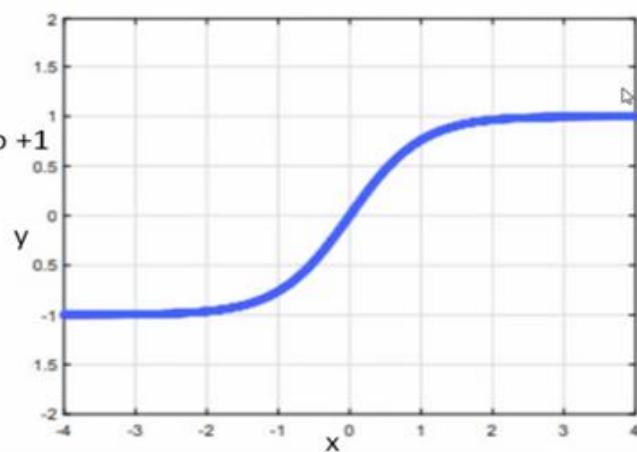
Value Range : 0 to 1



$$f(x) = \frac{1}{1 + e^{-x}}$$

### Tanh Function

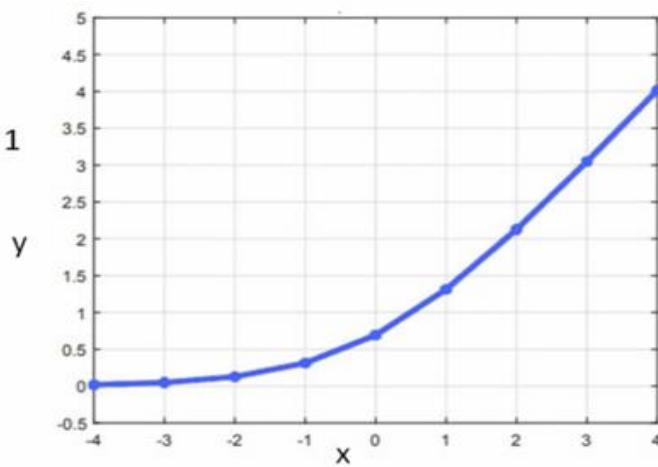
Value Range : -1 to +1



$$f(x) = \frac{1 - e^{-2x}}{1 + e^{-2x}}$$

### Softplus Function

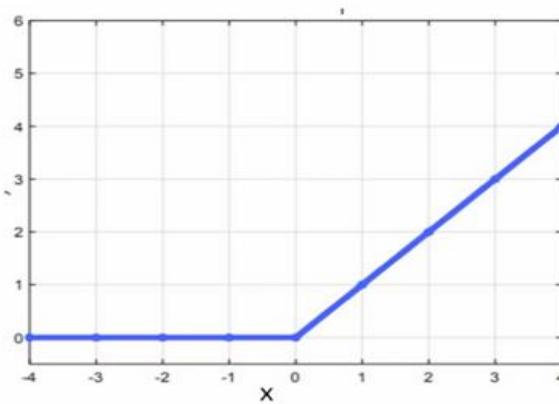
Value Range : 0 to 1



$$f(x) = \ln(1 + e^x)$$

### ReLU (Rectified linear unit) Function

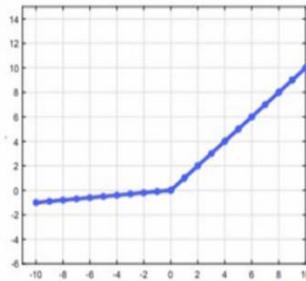
**Value Range : 0 to  $\infty$**



$$f(x) = \max(0, x)$$

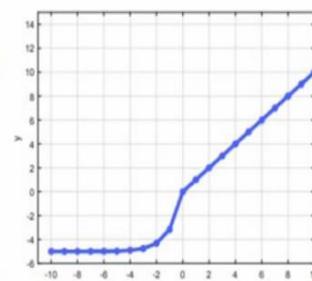
### Variations in ReLu Function

$$f(x) = \begin{cases} x & (x > 0) \\ a \cdot x & (\text{otherwise}) \end{cases}$$



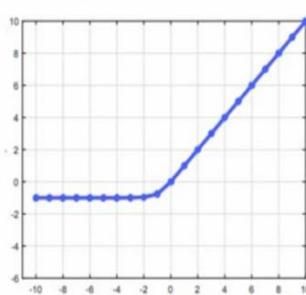
Leaky ReLU Function

$$f(x) = \begin{cases} x & (x \geq 0) \\ a \cdot (e^x - 1) & (\text{otherwise}) \end{cases}$$



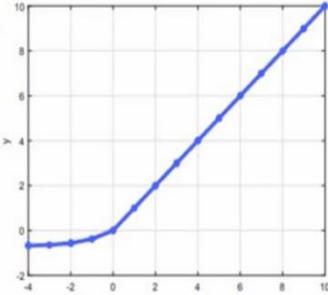
Elu Function

$$f(x) = \begin{cases} \frac{1-e^{-2x}}{1+e^{-2x}} & (x < 0) \\ \max(0, x) & (\text{otherwise}) \end{cases}$$



Tanh ReLU Function

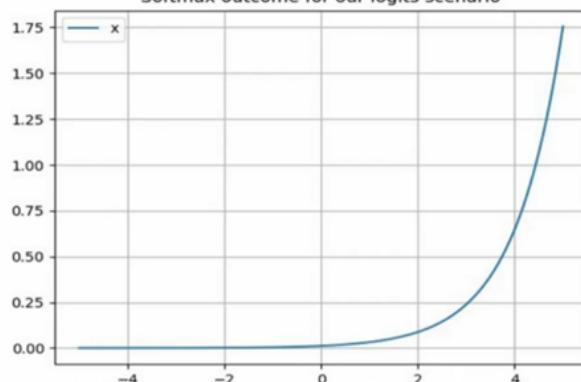
$$f(x) = \begin{cases} \ln(1 + e^x) - \ln 2 & (x < 0) \\ \max(0, x) & (\text{otherwise}) \end{cases}$$



Softplus (SmoothReLU) ReLU Function

### Softmax Function

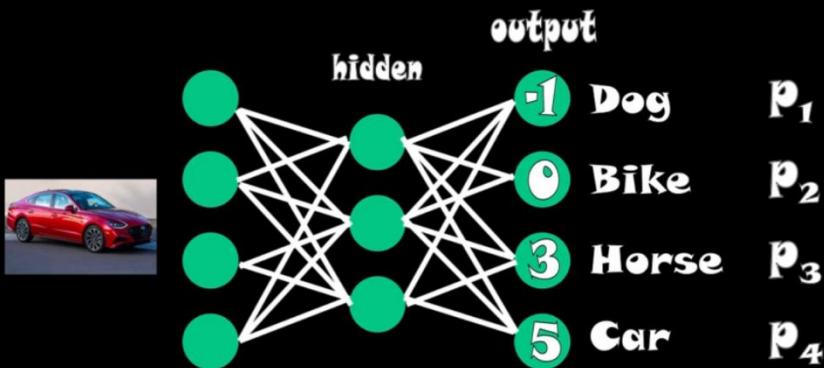
Softmax outcome for our logits scenario



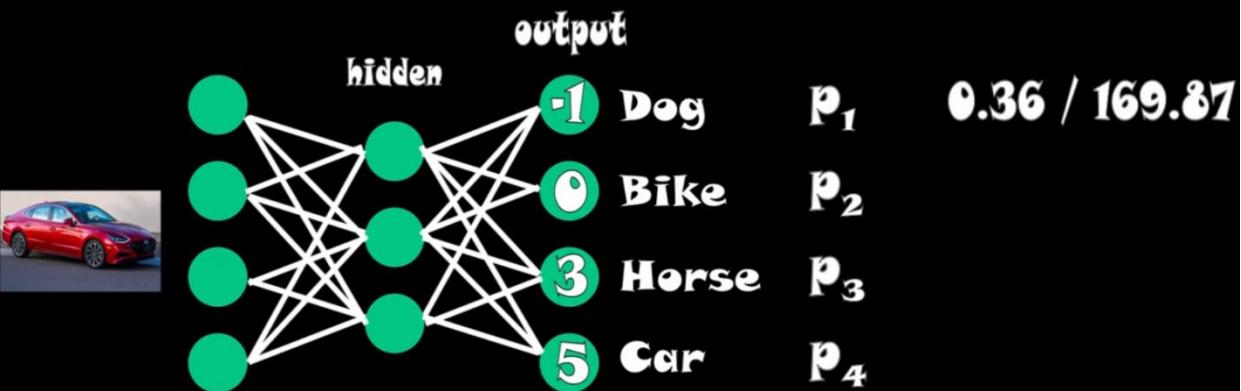
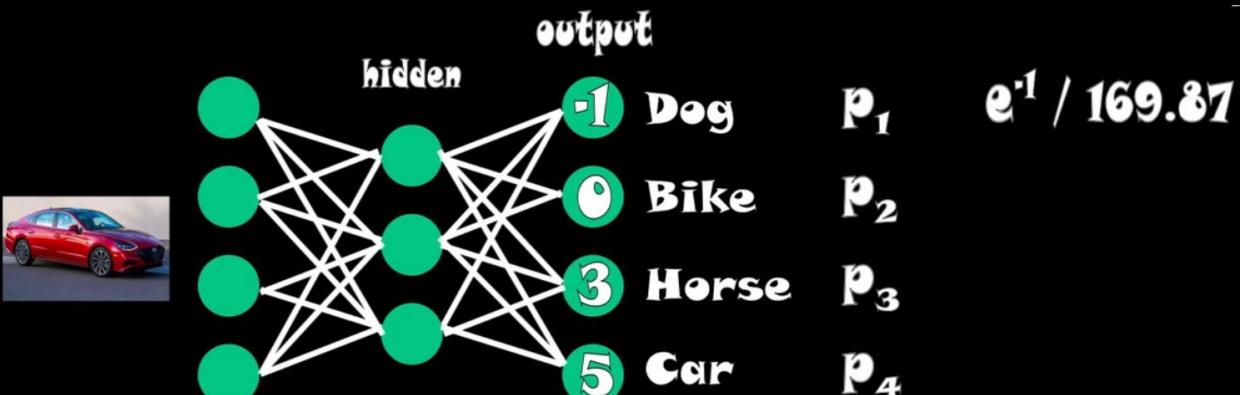
$$f_i(\vec{a}) = \frac{e^{a_i}}{\sum_k e^{a_k}}$$

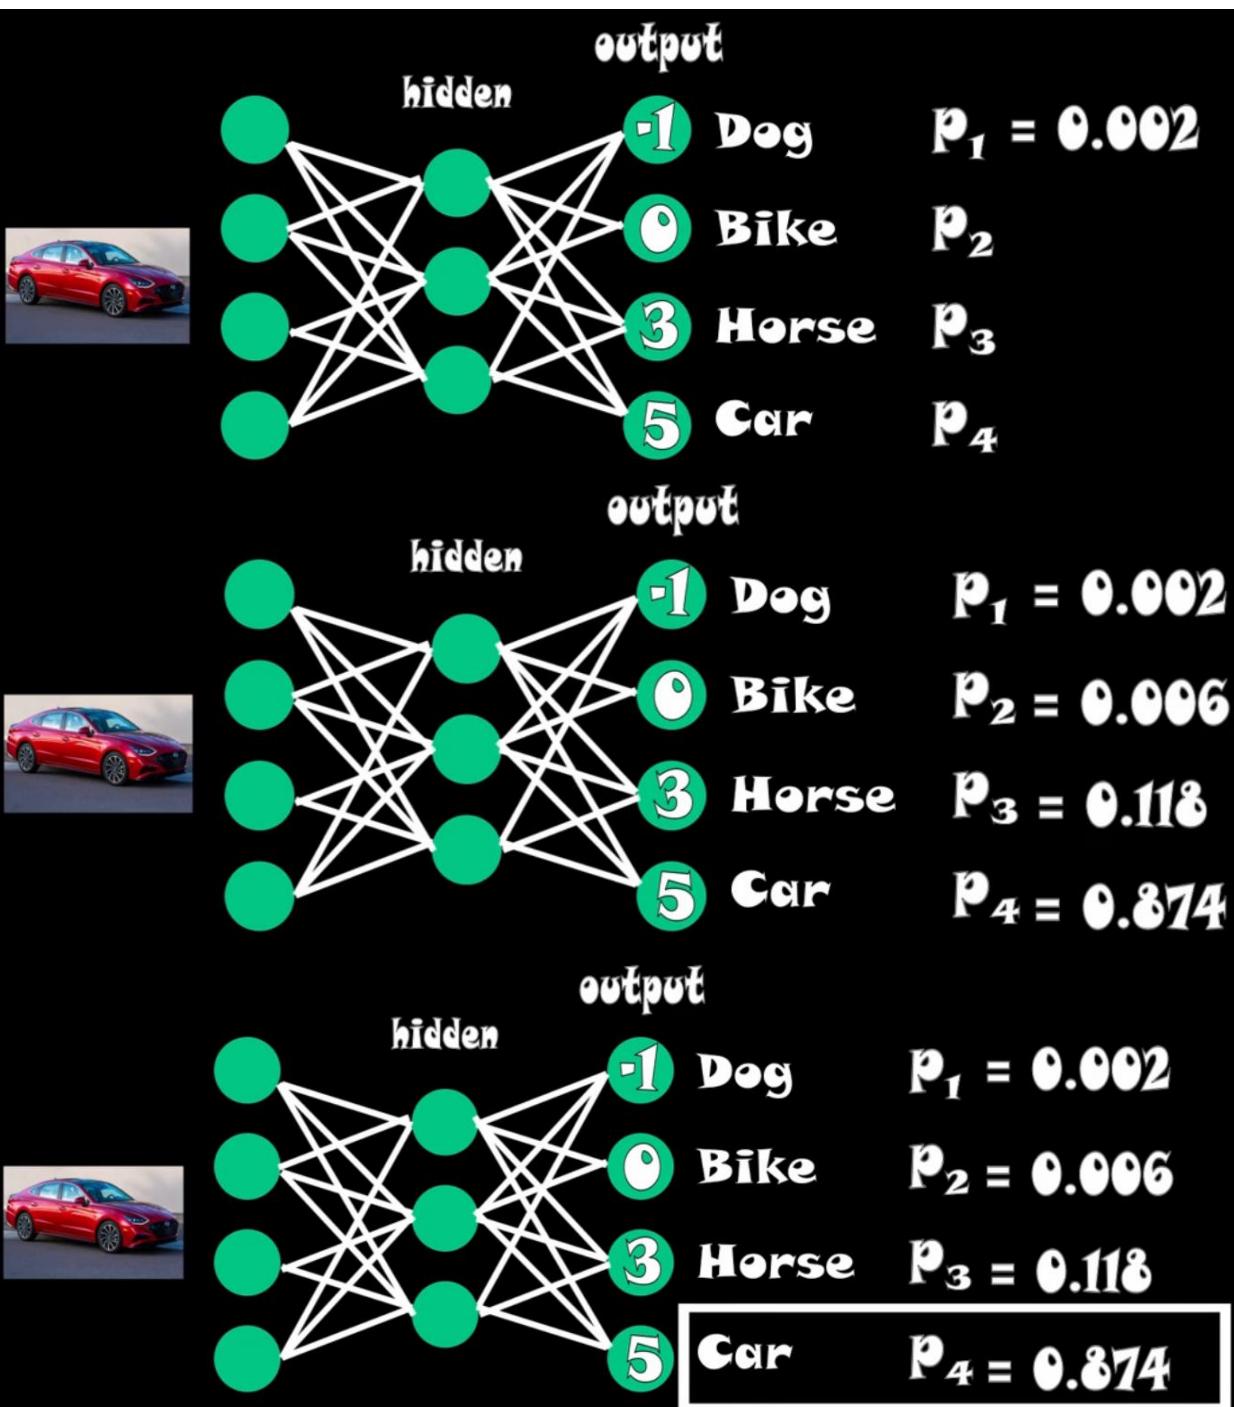
# Softmax

$$f(y_i) = \frac{e^{y_i}}{\sum_k e^{y_k}}$$



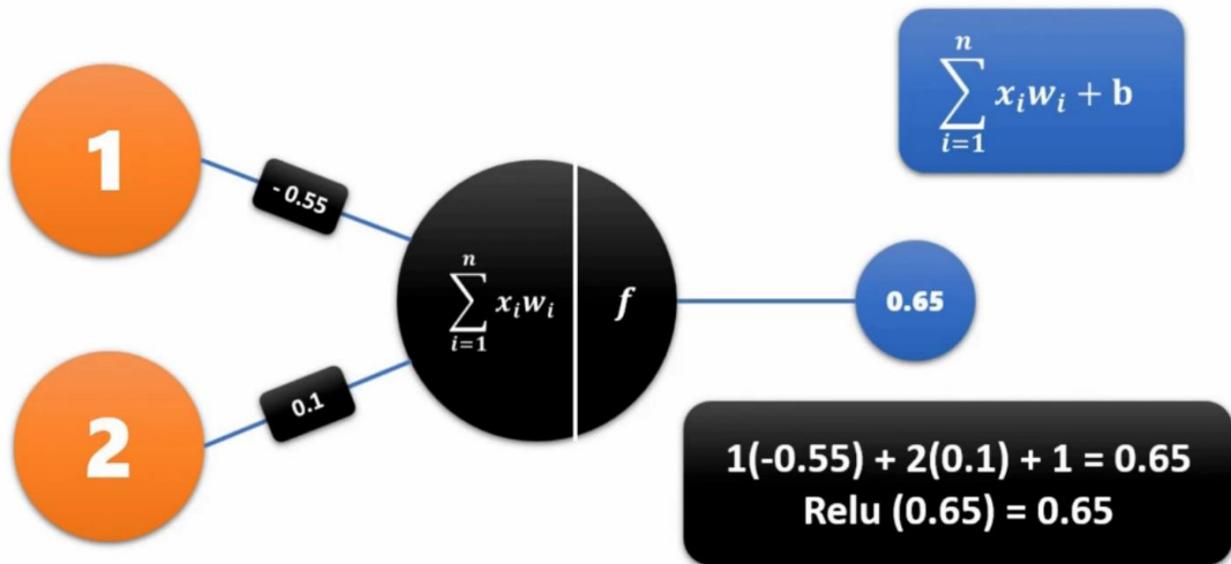
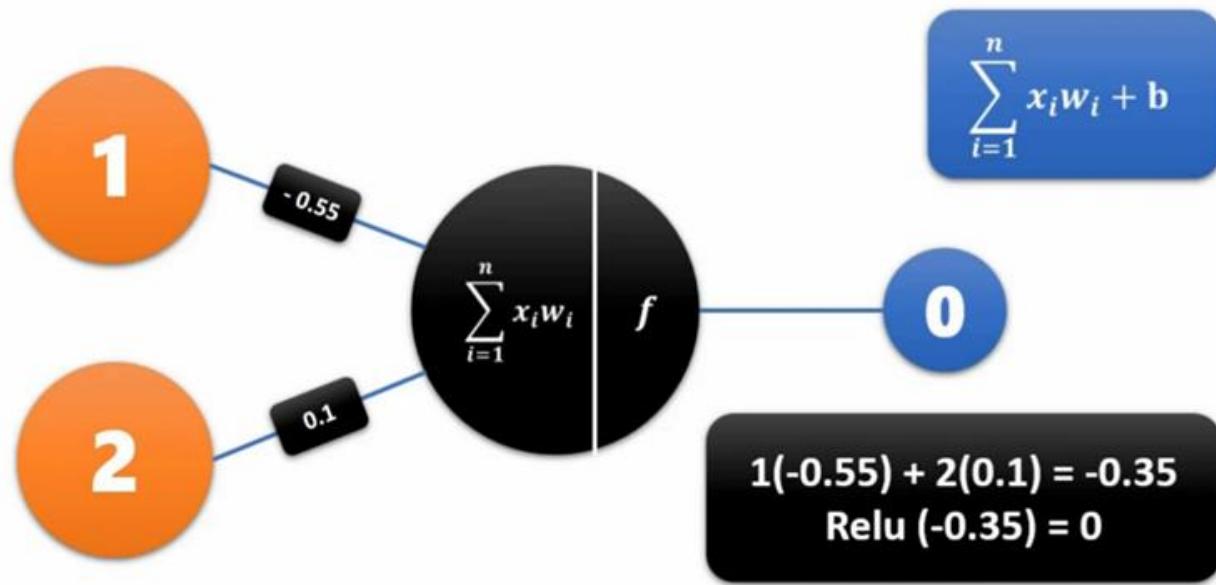
$$\sum_k e^{y_k} = e^{-1} + e^0 + e^3 + e^5 = 169.87$$





- 1 · Softmax function converts real values into probabilities.
- 2 · It only used as output layer of neural network.
- 3 · You can consider higher probability as actual output.

## What is the Role of Bias in ANN?



## Vanishing and Exploding Gradient

- Vanishing Gradient and Exploding Gradient Problem are difficulties found in training certain Artificial Neural Networks with gradient based methods like Back Propagation.
- In particular, this problem makes it really hard to learn and tune the parameters of the earlier layers in the network.

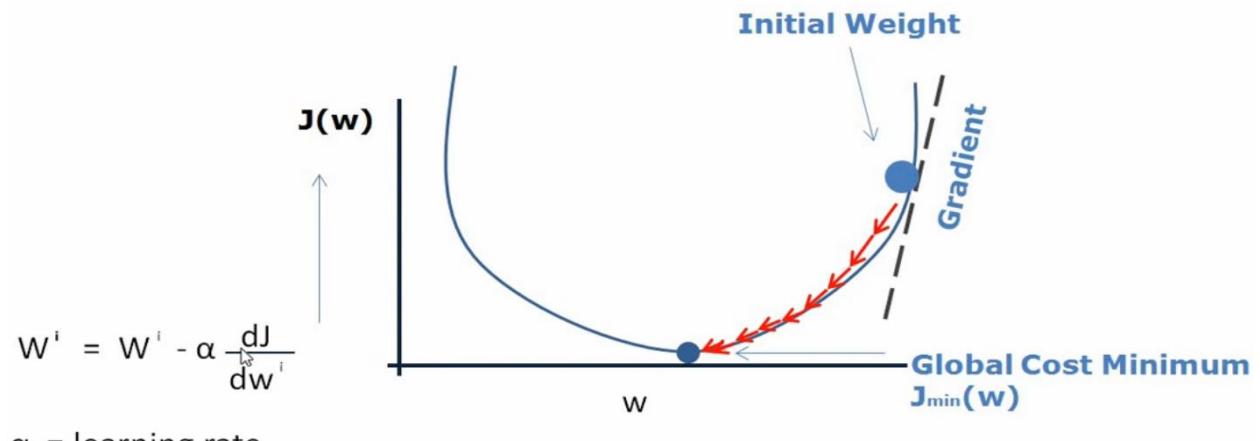
To calculate derivative of error w.r.t first weight, back-propagate via chain rule



$$\frac{dJ}{dw_1} = \frac{dJ}{d\hat{y}} * \frac{d\hat{y}}{da_2} * \frac{da_2}{da_1} * \frac{da_1}{dw_1}$$

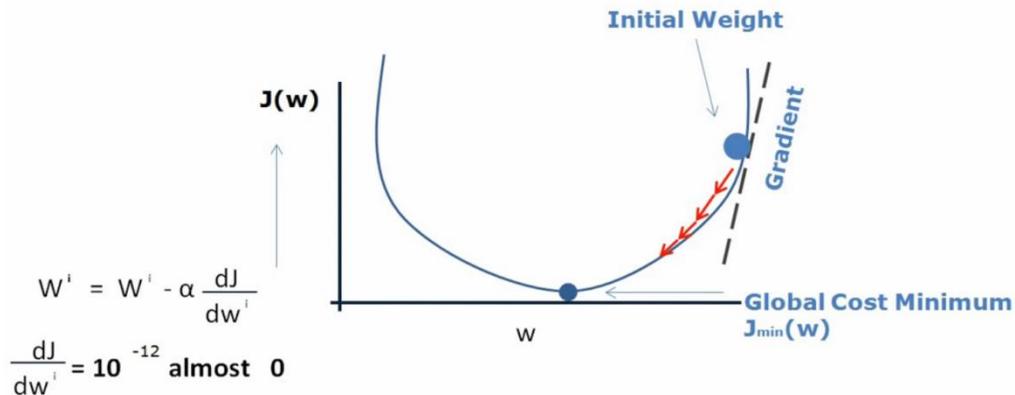
$$W_1 = W_1 - \alpha \frac{dJ}{dw_1}$$

**How does Gradient Descent work?**



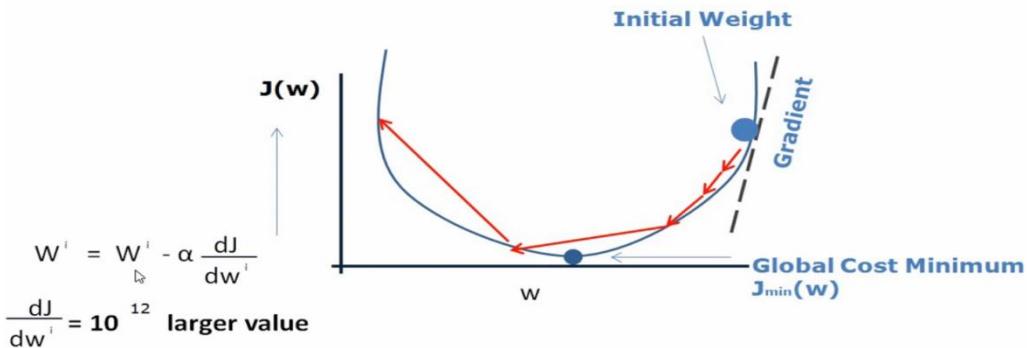
In Maths **Derivative** is a way to show rate of change at a given point

## Vanishing Gradient Problem



## Vanishing Gradient Problem

## Exploding Gradient Problem



## Exploding Gradient Problem

W ~ small -> Vanishing  
W ~ large -> Exploding

## Solution for Vanishing Gradient Problem

Vanishing Gradient Problem was because of Sigmoid function. Range 0-0.25  
Even if I use tanh same problem I will face Vanishing Gradient Problem tanh 0-1.

-The simplest solution is to use  
ReLU activation function instead of sigmoid or tanh activation  
which doesn't cause a small derivative.

-Use weight initializers.

-Use some other network

## Solution for Exploding Gradient Problem

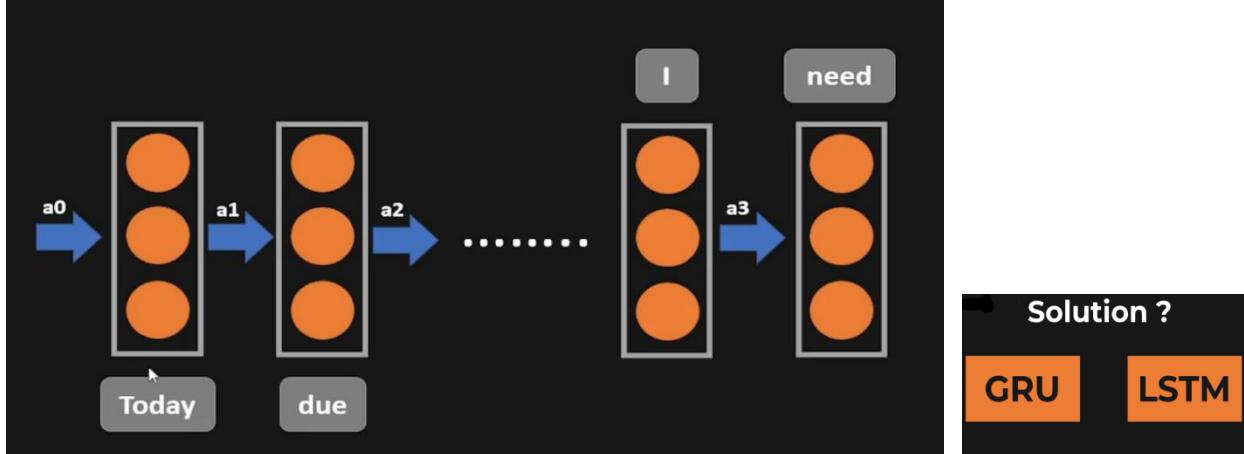
A solution to fix this is to apply gradient clipping.  
which places a predefined threshold on the gradients to prevent it  
from getting too large, and by doing this it doesn't change the direction  
of the gradients it only changes its length.

# VANISHING GRADIENT



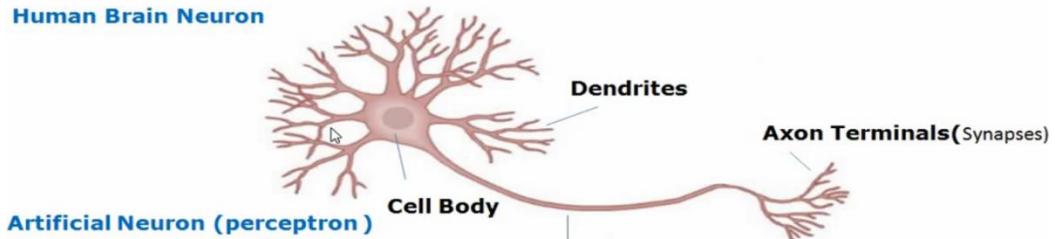
Today, due to my current job situation and family conditions, I **need** to take a loan.

Last Year, due to my current job situation and family conditions, I **had** to take a loan.

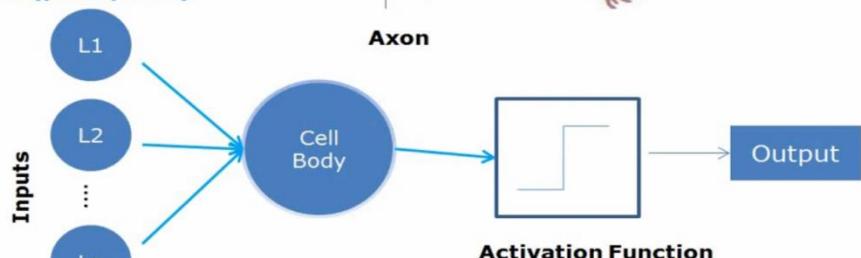


## Human Brain and Artificial Neuron

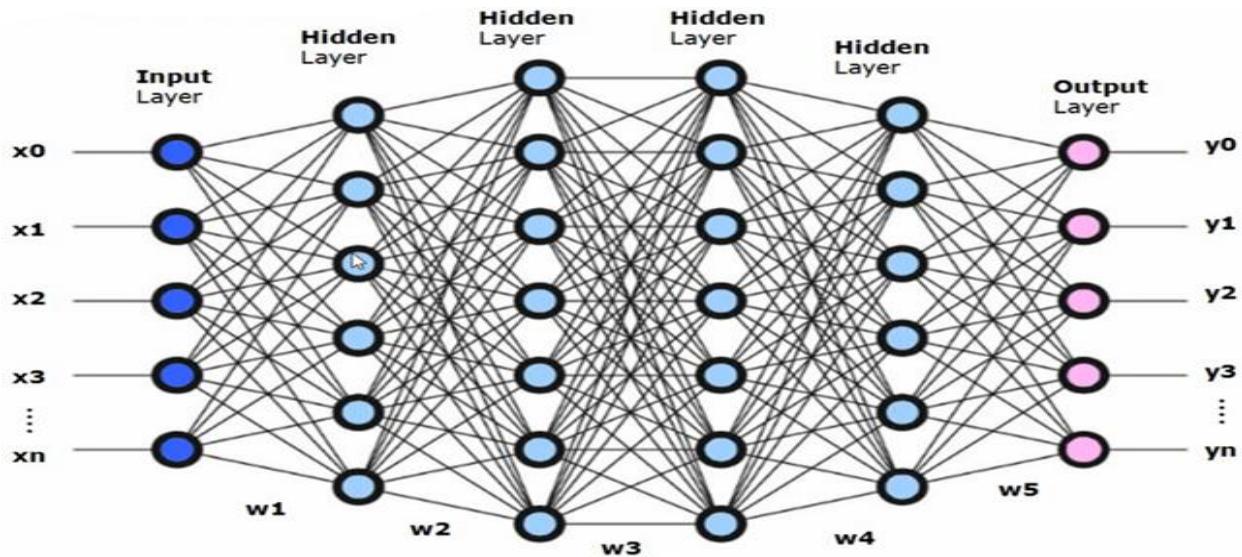
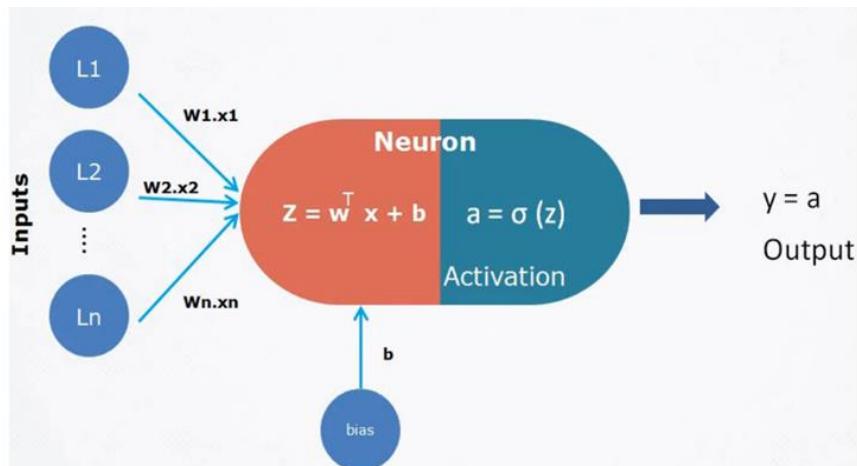
**Human Brain Neuron**



**Artificial Neuron (perceptron)**



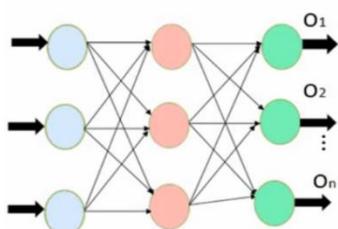
◀ ▶ ⌂ ⌃ ⌄



## Feed Forward Neural Network

- ✓ FFNN is one of the simplest forms of ANN
- ✓ A feedforward neural network is an artificial neural network wherein connections between the nodes do not form a cycle
- ✓ because inputs are processed only in the forward direction
- ✓ FFNN could be a Single layer perceptron or Multi layer perceptron
- ✓ We don't use it where order of sequence does matter. Like sentence as it predict only for current node , does not depend on previous.

### Mean squared error(MSE) in FFNN



$$E_{\text{total}} = \frac{1}{\text{Total output}} \sum (\text{Actual output} - \text{predicted output})^2$$
$$E_{\text{total}} = \frac{1}{n} \sum_{m=1}^n (o_m - \hat{o}_m)^2$$

Try to minimize the error for better prediction

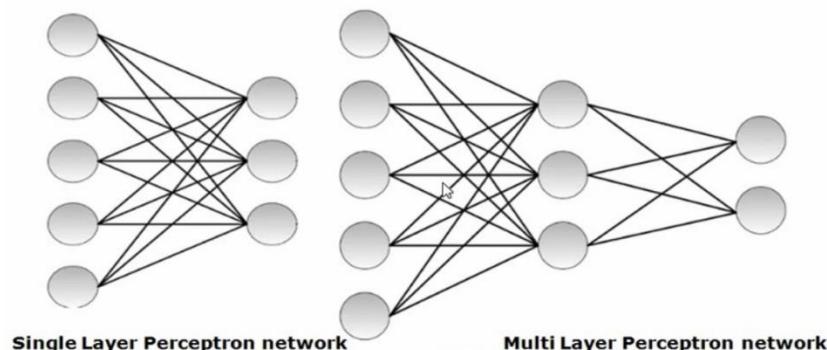
### Perceptron

#### Single Layer Perceptron (SLP) –

This is the simplest feedforward neural network and does not contain any hidden layer.

#### Multi Layer Perceptron (MLP)

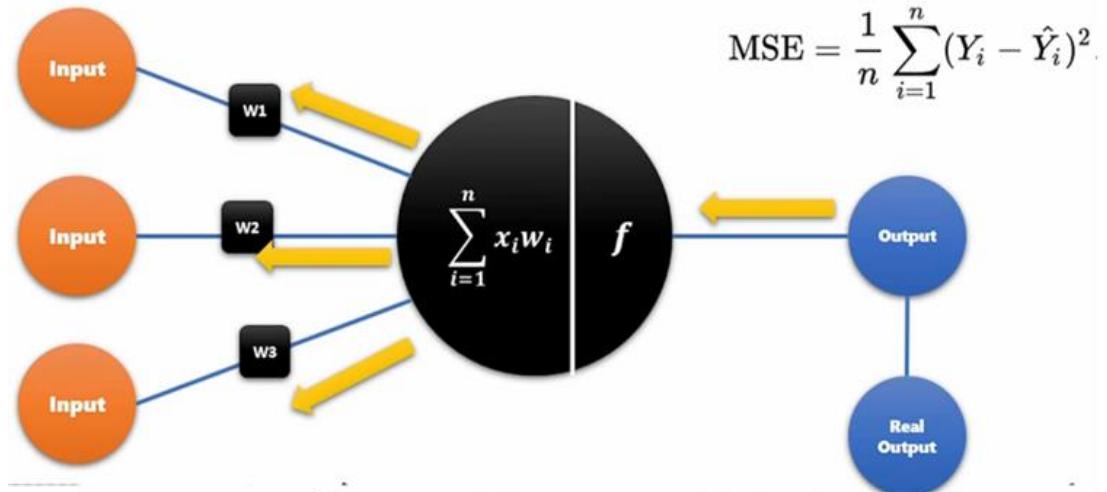
A Multi Layer Perceptron has one or more hidden layers.



### Applications

- ✓ Data Compression.
- ✓ Pattern Recognition.
- ✓ Computer Vision.
- ✓ Sonar Target Recognition.
- ✓ Speech Recognition.
- ✓ Handwritten Characters Recognition.

## Back Propagation Neural Network



**Try to minimize the error for better prediction model**

**Error Calculation** – difference of model output from the actual output.

**Minimum Error** – Check whether the error is minimized or not.

**Update the parameters** – If the error is huge then, update the parameters weights and biases and again check the error. Repeat the process until the error becomes minimum.



**Model is ready to make a prediction** – Now model will produce the output more accurate.

## Applications of Backpropagation

- i. **Speech Recognition**
- ii. **Face Recognition**
- iii. **Signature Recognition**

## Convolutional Neural Network (CNN)

CNN is a special type of Feed Forward Neural network

- ✓ Images recognition
- ✓ Images classifications
- ✓ Objects detections
- ✓ Faces Recognition

**Convolutional Neural Network (CNN)**

Filters/Kernals		
-1	1	-1
-1	1	-1
1	1	1
1	1	1
-1	1	-1
-1	1	-1
-1	1	-1
1	1	1
-1	1	-1
-1	1	-1
-1	1	-1

Training Data				
-1	-1	1	-1	-1
-1	-1	1	-1	-1
1	1	1	1	1
-1	-1	1	-1	-1
-1	-1	1	-1	-1

Testing Data				
-1	-1	1	-1	-1
-1	-1	1	-1	-1
1	1	1	1	1
-1	-1	1	-1	-1
-1	-1	1	-1	-1

KotaCoaching.com

**Filters**

-1	1	-1
-1	1	-1
1	1	1
1	1	1
-1	1	-1
-1	1	-1
-1	1	-1

$(1+1+1+1+1+1+1+1+1)/9 = 1$

-1	-1	1	-1	-1
-1	1	1	-1	-1
1	1	1	1	1
-1	-1	1	-1	-1
-1	-1	1	-1	-1

1	1	1
1	1	1
1	1	1

**Stride** is the number of pixels shifts over the input matrix.  
Here Stride = 1

KotaCoaching.com

Filters

-1	1	-1
-1	1	-1
1	1	1

1	1	1
-1	1	-1
-1	1	-1

-1	1	-1
1	1	1
-1	1	-1

-1	-1	1	-1	-1
-1	-1	1	-1	-1
1	1	1	1	1
-1	-1	1	-1	-1
-1	-1	1	-1	-1

to make the final filter more informative  
we use padding in image matrix

Featured Images

.11	1	.33
-.33	.11	-.33
-.33	.11	-.11

-.33	.11	-.33
-.33	.11	-.33
.11	1	-.11

-.11	.11	-.33
-33	1	.33
-33	.11	-.11

KotaCoaching.com

## ReLU Layer

Filters

-1	1	-1
-1	1	-1
1	1	1

1	1	1
-1	1	-1
-1	1	-1

-1	1	-1
1	1	1
-1	1	-1

-1	-1	1	-1	-1
-1	-1	1	-1	-1
1	1	1	1	1
-1	-1	1	-1	-1
-1	-1	1	-1	-1

$$f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

KotaCoaching.com

Featured Images

.11	1	.33
0	.11	0
0	.11	0

0	.11	0
0	.11	0
.11	1	0

0	.11	0
0	1	.33
0	.11	0

ReLu Activation

Filters

-1	1	-1
-1	1	-1
1	1	1

1	1	1
-1	1	-1
-1	1	-1

-1	1	-1
1	1	1
-1	1	-1

## Pooling Layer

Featured Images

.11	1	.33
0	.11	0
0	.11	0

1		

0	.11	0
0	.11	0
.11	1	0

0	.11	0
0	1	.33
0	.11	0

Pooling can be of different types

- Max Pooling
- Average Pooling
- Sum Pooling

KotaCoaching.com

Activate Windows

Filters

-1	1	-1
-1	1	-1
1	1	1

1	1	1
-1	1	-1
-1	1	-1

-1	1	-1
1	1	1
-1	1	-1

## Pooling Layer

Featured Images

.11	1	.33
0	.11	0
0	.11	0

1	1
.11	.11

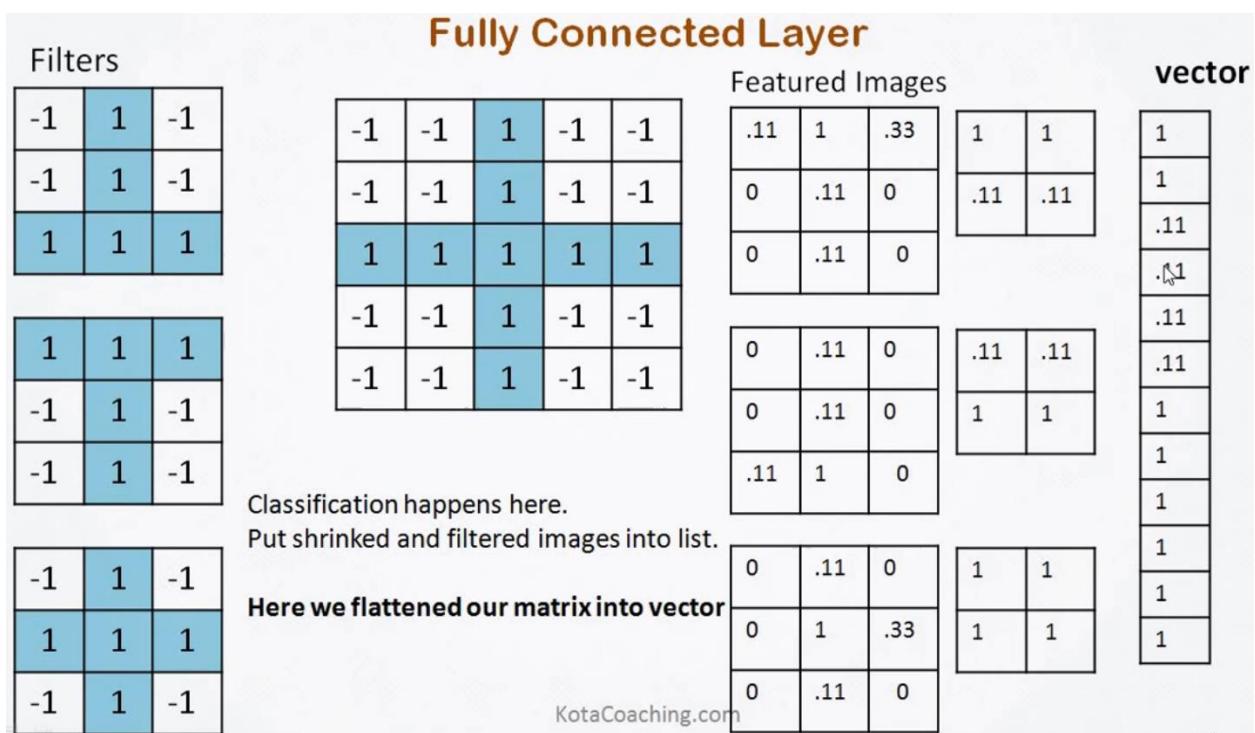
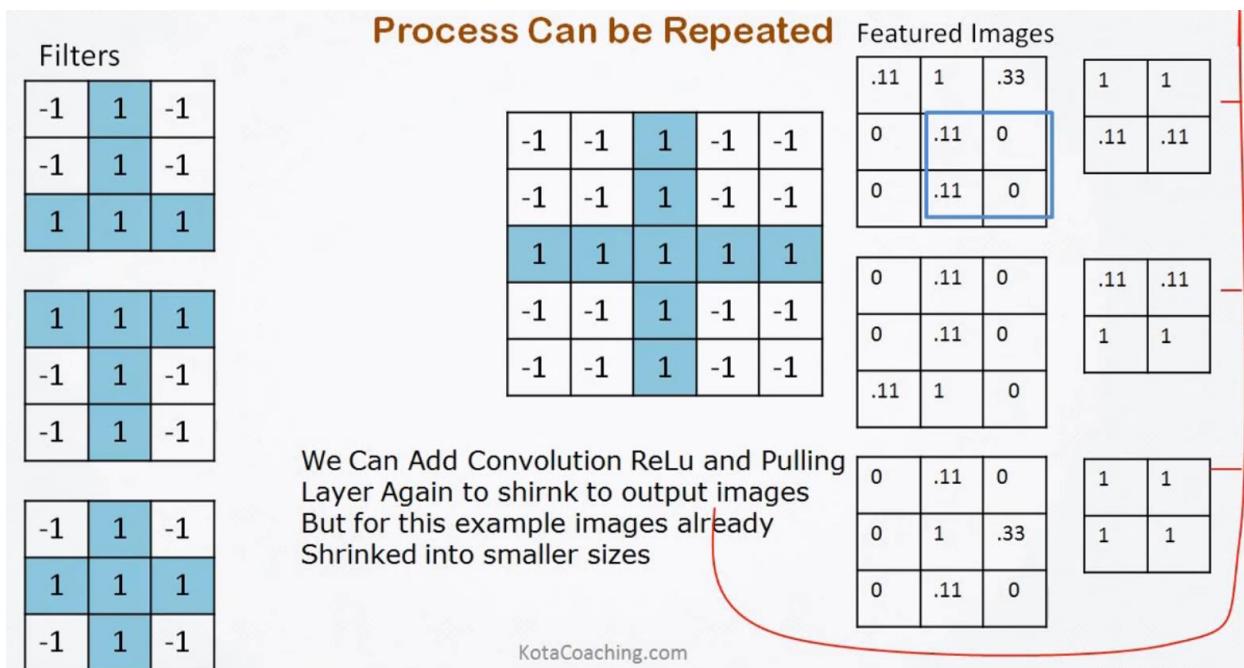
0	.11	0
0	.11	0
.11	1	0

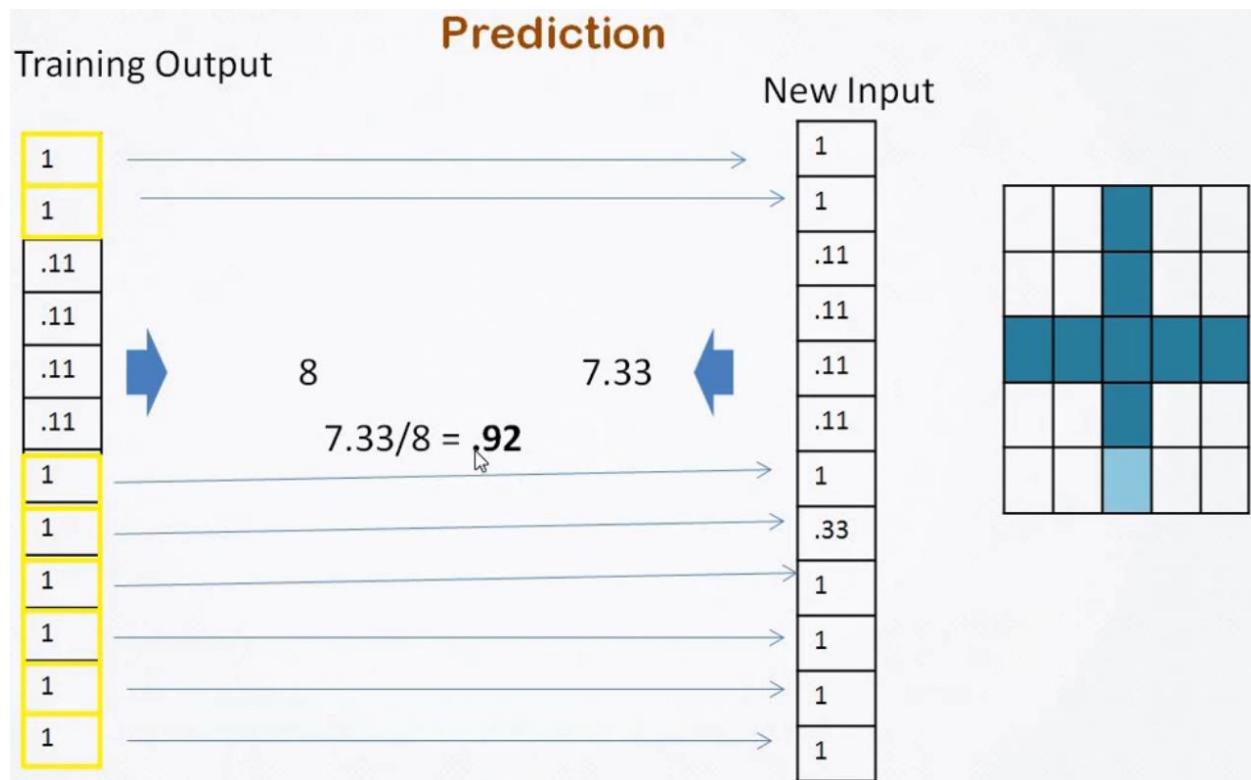
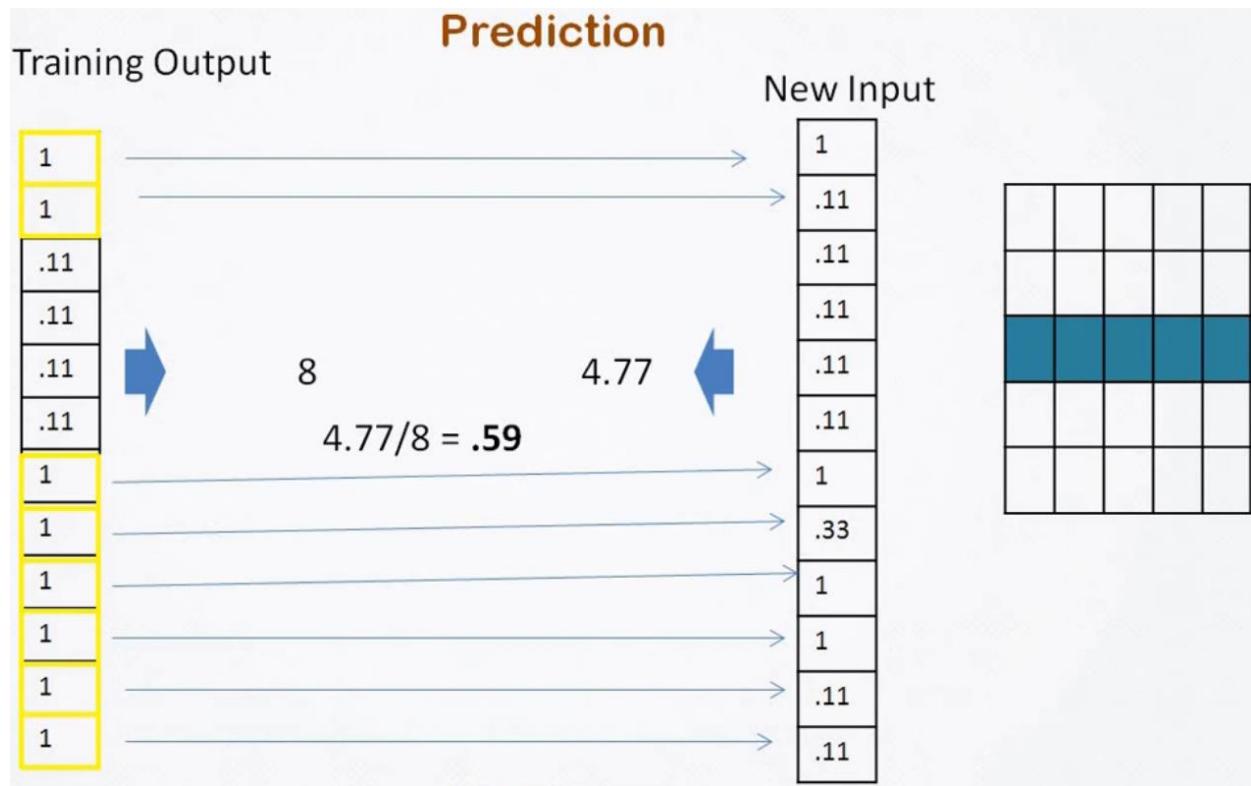
.11	.11
1	1

Shrink image in smaller size.  
Move window size 2 across filtered image.  
Pick highest value

KotaCoaching.com

Activate Windows





0	0	0	0	0	0	0
0	-1	-1	1	-1	-1	0
0	-1	-1	1	-1	-1	0
0	1	1	1	1	1	0
0	-1	-1	1	-1	-1	0
0	-1	-1	1	-1	-1	0
0	0	0	0	0	0	0

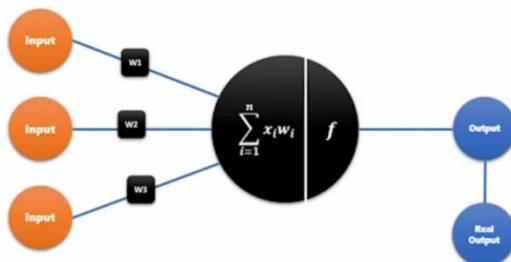
**Padding** : it refers to the amount of pixels added to an image when it is being processed by the kernel of a CNN  
 Padding is simply a process of adding layers of zeros to our input image

## Computer Vision

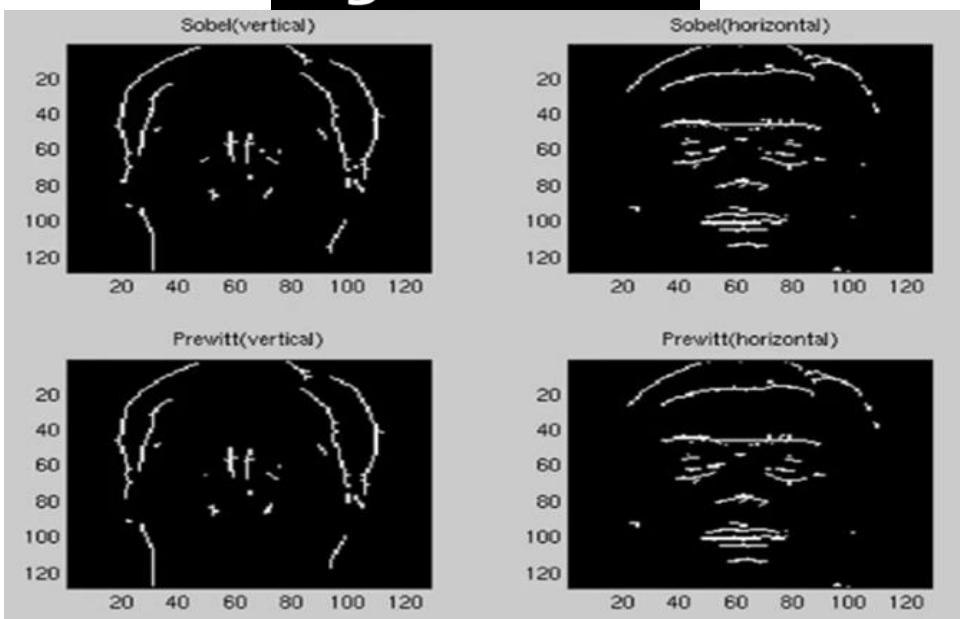


1920 x 1080 x 3

6,220,800 px



## Edge Detection



## Vertical Edge Detection

$6 \times 6$

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

Filter

$3 \times 3$

1	0	-1
1	0	-1
1	0	-1

\*

=

$4 \times 4$

-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

$$(3 \times 1) + (1 \times 1) + (2 \times 1) + (0 \times 0) + (5 \times 0) + (7 \times 0) + (1 \times -1) + (8 \times -1) + (2 \times -1) = -5$$

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



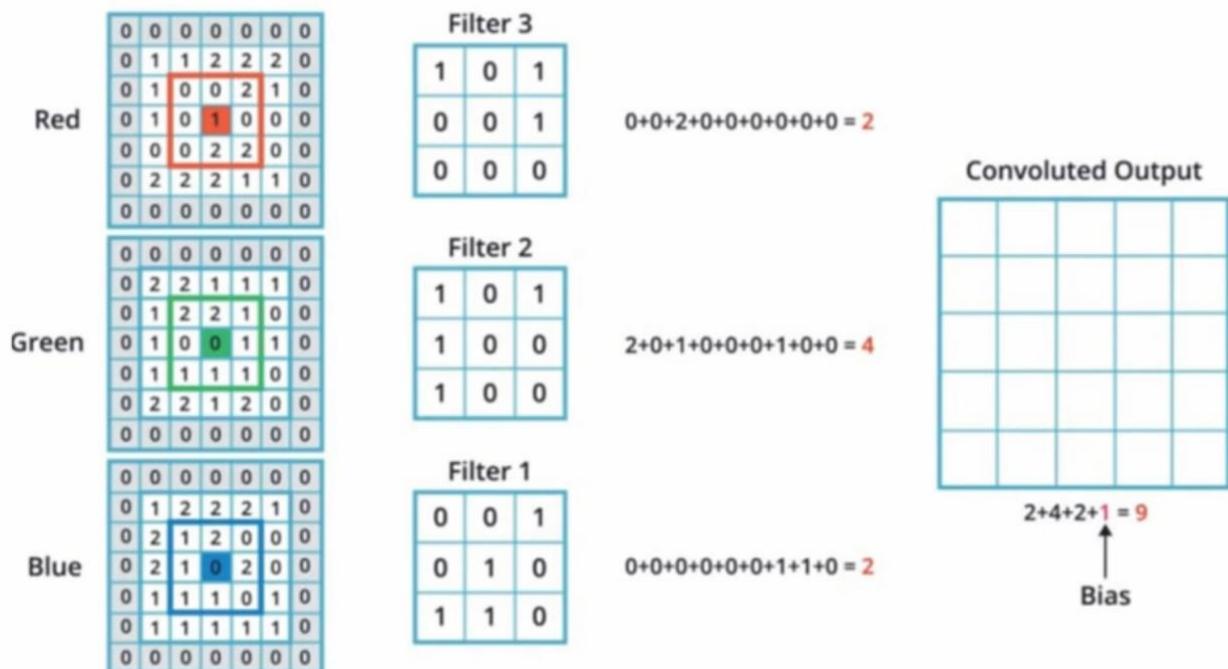
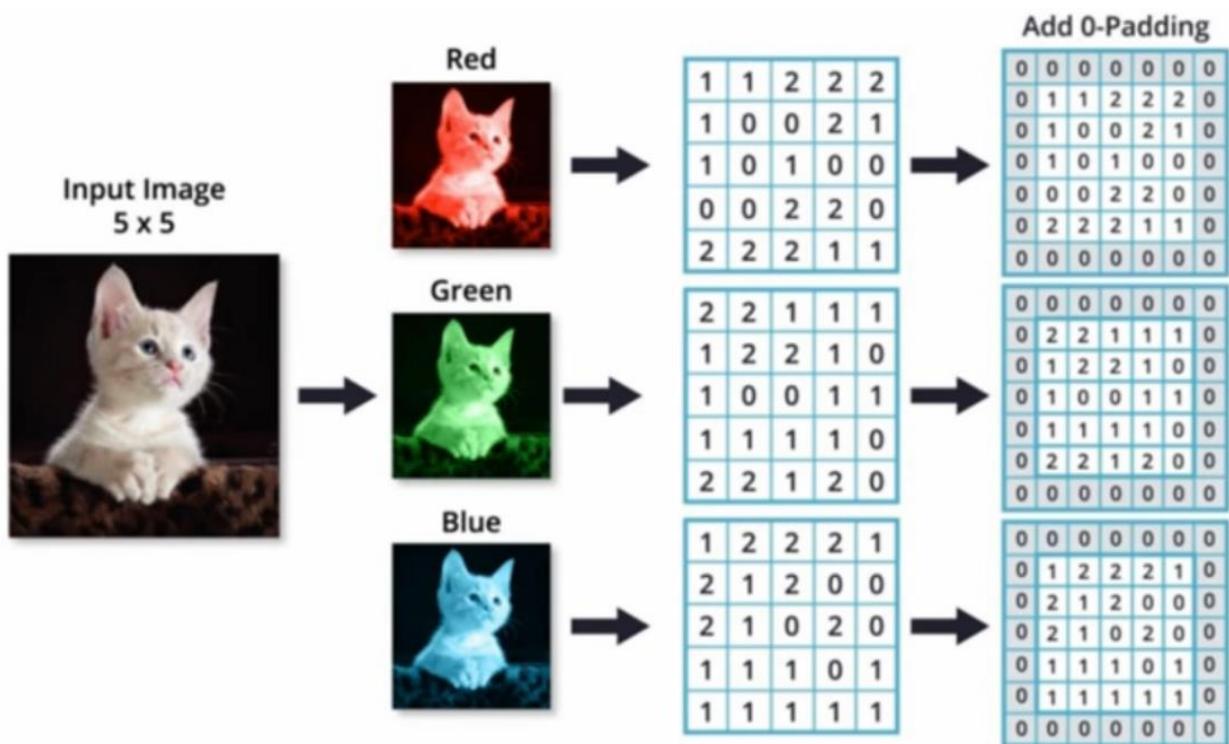
\*

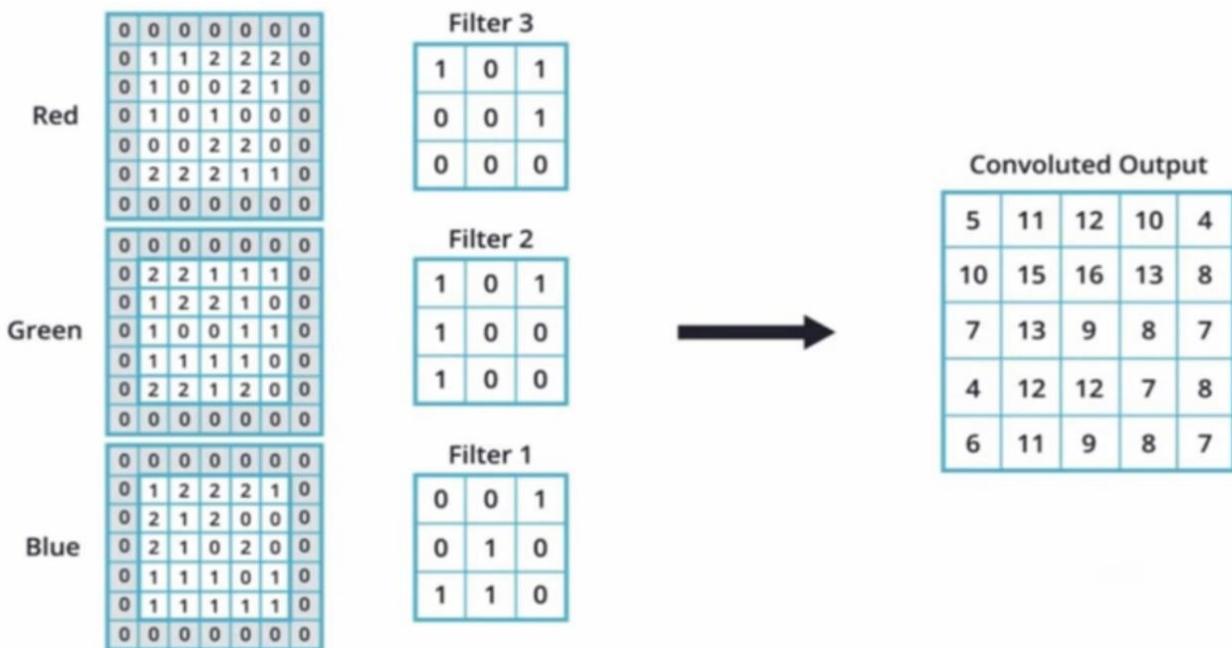
1	0	-1
1	0	-1
1	0	-1

=

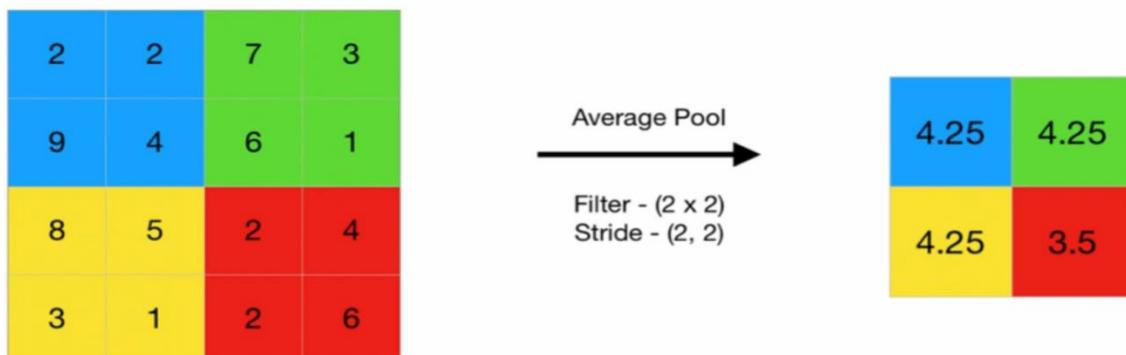
0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



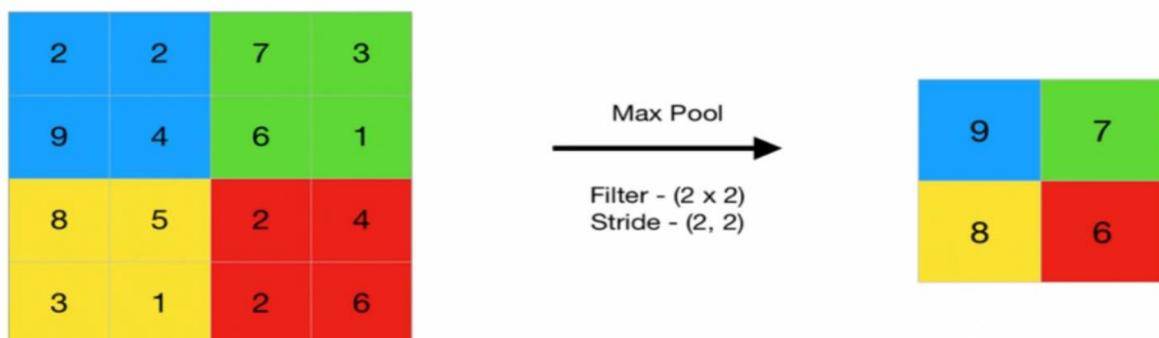




## Average Pooling



## Max Pooling



## ▼ Importing Libraries and Image

```
✓ 0s  [1] import numpy as np
    from scipy import misc
    i = misc.face(gray=True)

✓ 1s  [2] import matplotlib.pyplot as plt
    plt.grid(False)
    plt.gray()
    plt.axis('off')
    plt.imshow(i)
    plt.show()
```



```
[ ] i_transformed = np.copy(i)
size_x = i_transformed.shape[0]
size_y = i_transformed.shape[1]
```

## ▼ Apply Filter On image

```
[ ] #Experiment with different values for fun effects.
#filter = [ [0, 1, 0], [1, -4, 1], [0, 1, 0]]

# A couple more filters to try for fun!
#filter = [ [-1, -2, -1], [0, 0, 0], [1, 2, 1]]
#filter = [ [-1, 0, 1], [-2, 0, 2], [-1, 0, 1]]

filter = [[1,0,-1],
          [1,0,-1],
          [1,0,-1]]
```

$$\begin{array}{c}
 \text{6} \times 6 \\
 \begin{array}{|c|c|c|c|c|c|} \hline
 3 & 0 & 1 & 2 & 7 & 4 \\ \hline
 1 & 5 & 8 & 9 & 3 & 1 \\ \hline
 2 & 7 & 2 & 5 & 1 & 3 \\ \hline
 0 & 1 & 3 & 1 & 7 & 8 \\ \hline
 4 & 2 & 1 & 6 & 2 & 8 \\ \hline
 2 & 4 & 5 & 2 & 3 & 9 \\ \hline
 \end{array}
 \end{array}
 \quad *
 \begin{array}{c}
 \text{Filter} \\
 3 \times 3 \\
 \begin{array}{|c|c|c|} \hline
 1 & 0 & -1 \\ \hline
 1 & 0 & -1 \\ \hline
 1 & 0 & -1 \\ \hline
 \end{array}
 \end{array}
 = 
 \begin{array}{c}
 \text{4} \times 4 \\
 \begin{array}{|c|c|c|c|} \hline
 -5 & -4 & 0 & 8 \\ \hline
 -10 & -2 & 2 & 3 \\ \hline
 0 & -2 & -4 & -7 \\ \hline
 -3 & -2 & -3 & -16 \\ \hline
 \end{array}
 \end{array}$$

```

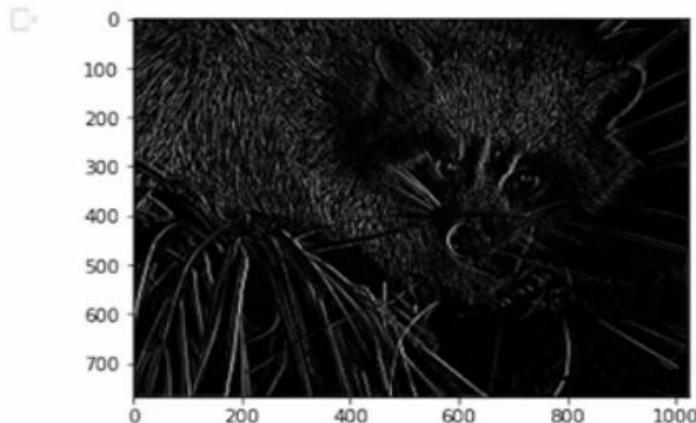
22s ✓ [8] for x in range(1,size_x-1):
    for y in range(1,size_y-1):
        convolution = 0.0
        convolution = convolution + (i[x-1, y-1] * filter[0][0])
        convolution = convolution + (i[x-1, y] * filter[0][1])
        convolution = convolution + (i[x-1, y+1] * filter[0][2])
        convolution = convolution + (i[x, y-1] * filter[1][0])
        convolution = convolution + (i[x, y] * filter[1][1])
        convolution = convolution + (i[x, y+1] * filter[1][2])
        convolution = convolution + (i[x+1, y-1] * filter[2][0])
        convolution = convolution + (i[x+1, y] * filter[2][1])
        convolution = convolution + (i[x+1, y+1] * filter[2][2])
        if(convolution<0):
            convolution=0
        if(convolution>255):
            convolution=255
        i_transformed[x, y] = convolution

```

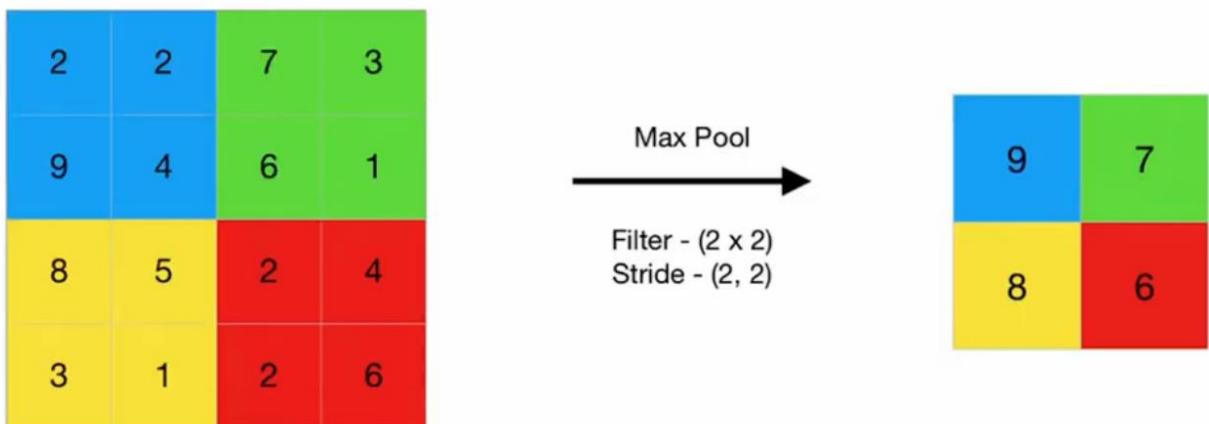
```

1s ✓ [8] # Plot the image. Note the size of the axes -- they are 512 by 512
plt.gray()
plt.grid(False)
plt.imshow(i_transformed)
#plt.axis('off')
plt.show()

```

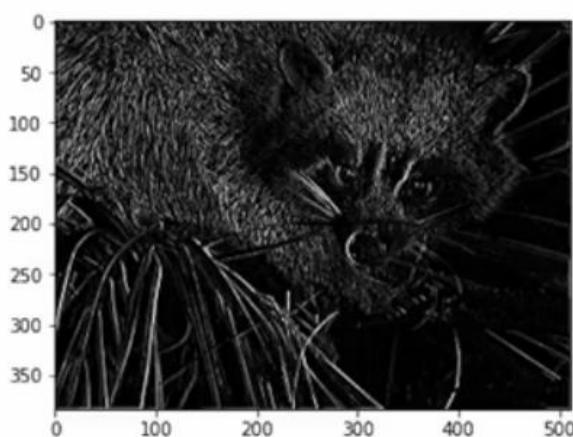


## Max Pooling



```
new_x = int(size_x/2)
new_y = int(size_y/2)
newImage = np.zeros((new_x, new_y))
for x in range(0, size_x, 2):
    for y in range(0, size_y, 2):
        pixels = []
        pixels.append(i_transformed[x, y])
        pixels.append(i_transformed[x+1, y])
        pixels.append(i_transformed[x, y+1])
        pixels.append(i_transformed[x+1, y+1])
        newImage[int(x/2),int(y/2)] = max(pixels)

# Plot the image. Note the size of the axes -- now 256 pixels instead of 512
plt.gray()
plt.grid(False)
plt.imshow(newImage)
# plt.axis('off')
plt.show()
```

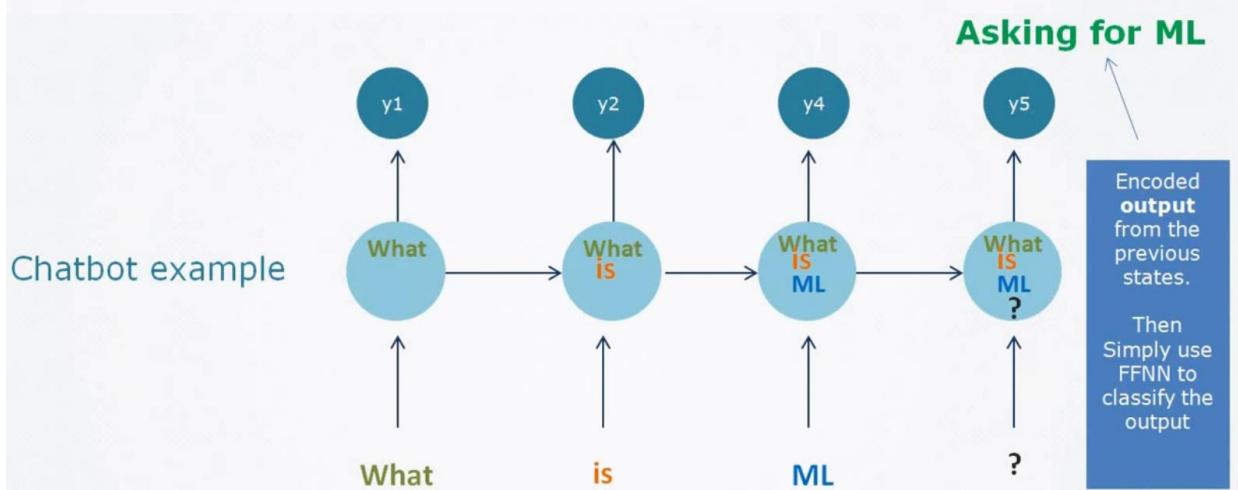


## Recurrent Neural Networks (RNN)

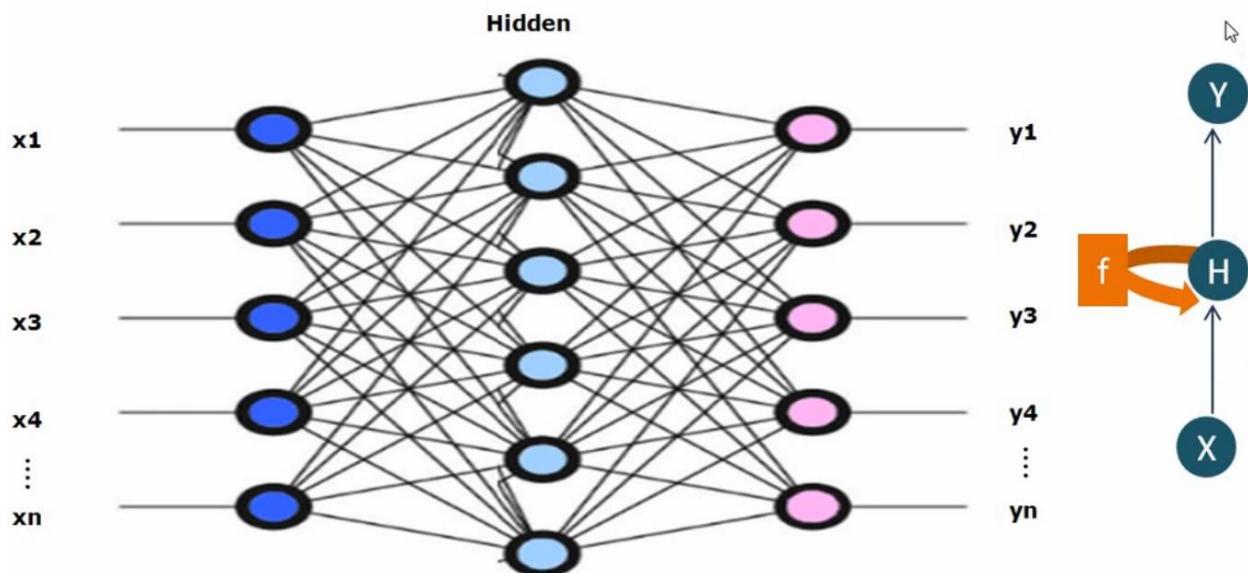
- Recurrent Neural Networks(**RNN**) are a type of Neural Network where the output from the previous step is fed as input to the current step
- *RNNs is that they have a "memory" which captures information about what has been calculated so far.*
- **RNN's** are mainly used for, Sequence Classification — Sentiment Classification

## Sequential Data

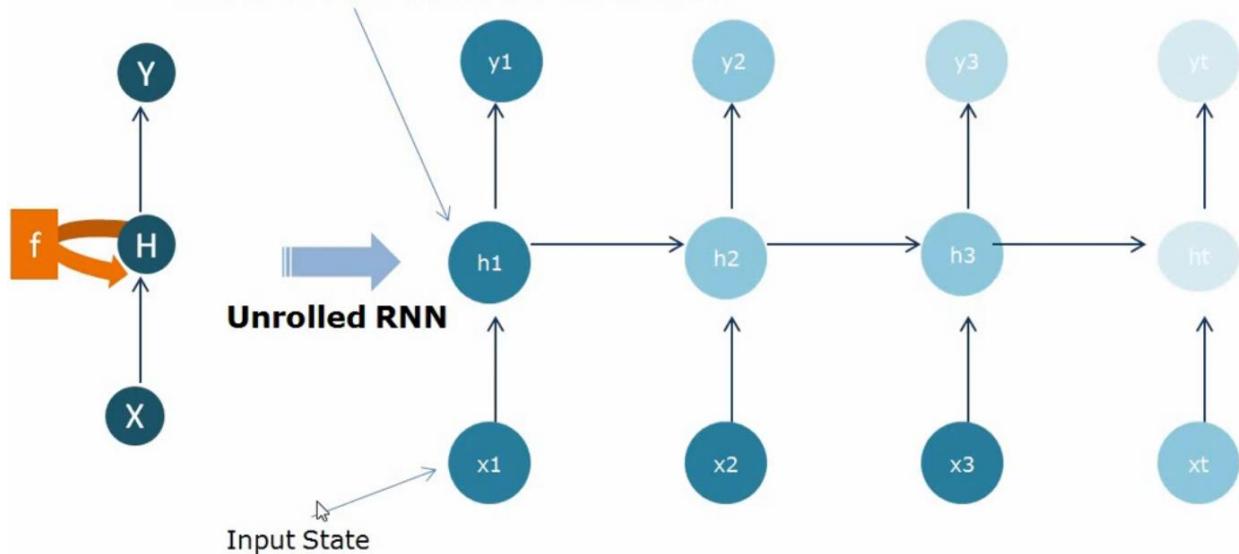
- Audio
- Video
- Text
- Financial Data

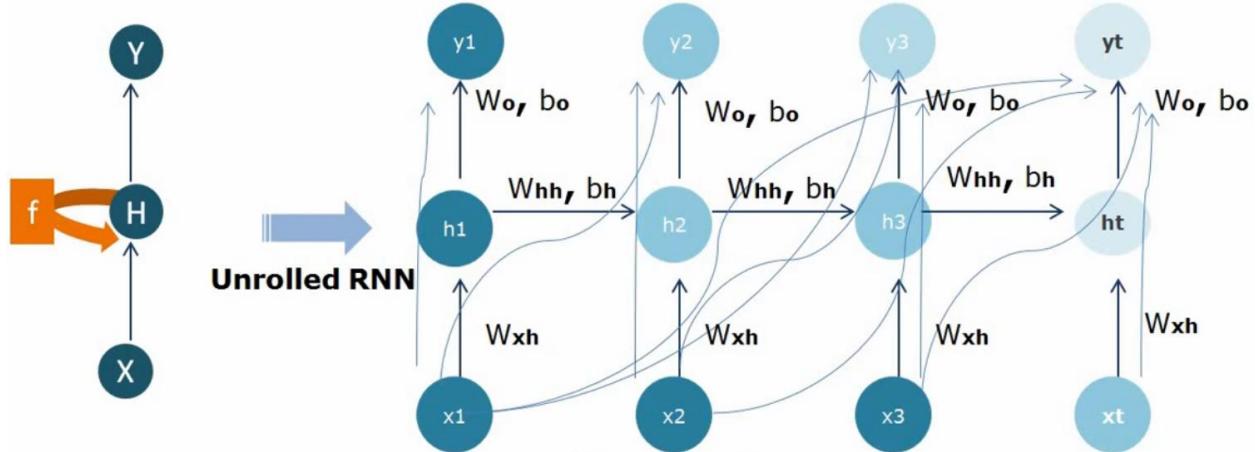


Make the hidden state depend on the previous hidden state



Hidden State  
hidden state- representation of previous information





$$h_t = \sigma(W_{xh}^T x_t + W_{hh}^T h_{t-1} + b_h)$$

$\tanh$  activation function

$$\hat{y}_t = \sigma(W_o^T h_t + b_o)$$

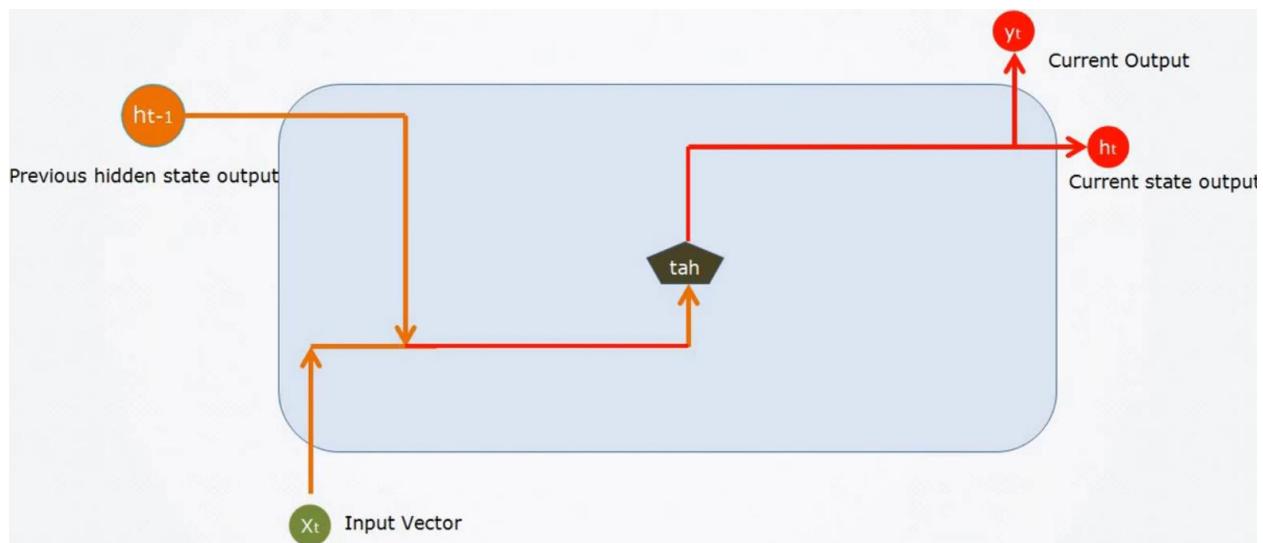
$$h_1 = \sigma(W_{xh}^T x_1 + W_{hh}^T h_0 + b_h)$$

$$\hat{y}_1 = \sigma(W_o^T h_1 + b_o)$$

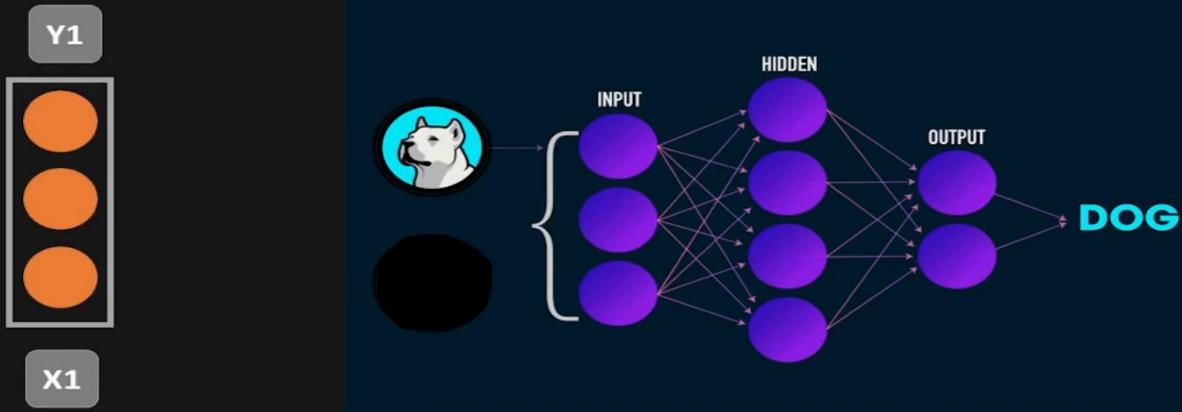
you have to initialize an  $h_0$  at the first hidden state.

$$h_2 = \sigma(W_{xh}^T x_2 + W_{hh}^T h_1 + b_h)$$

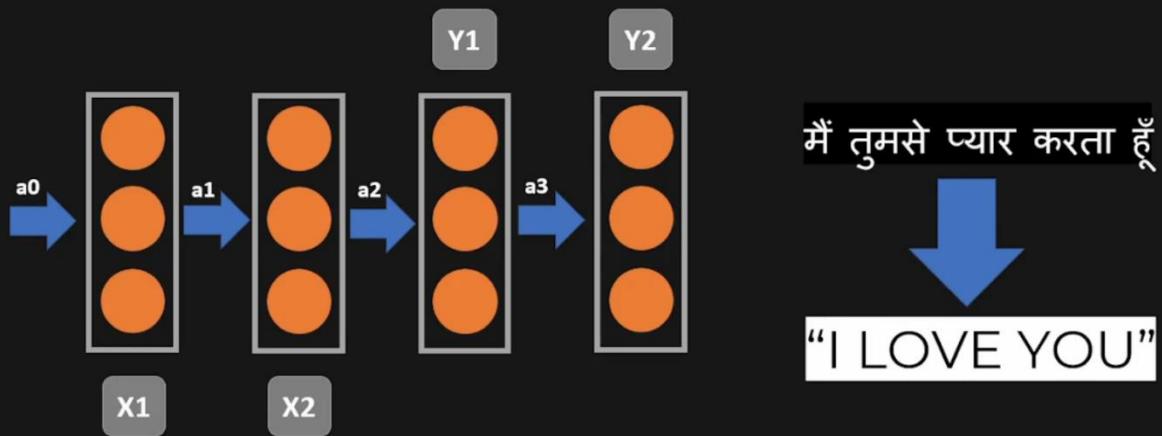
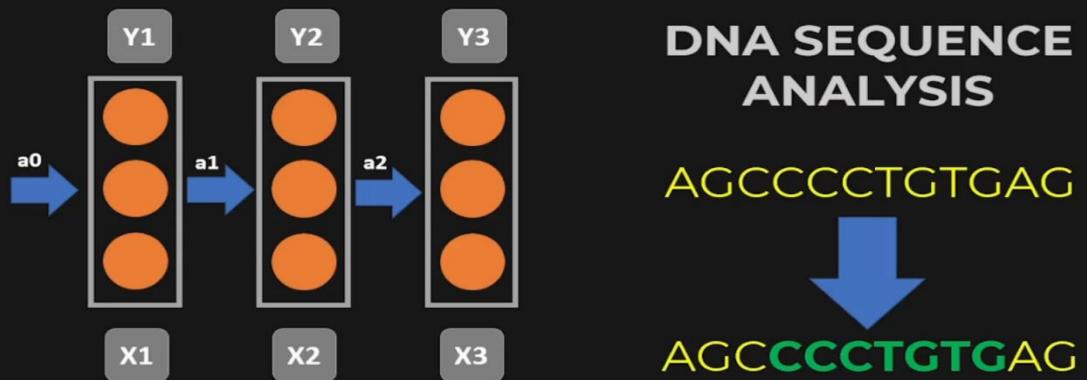
$$\hat{y}_2 = \sigma(W_o^T h_2 + b_o)$$



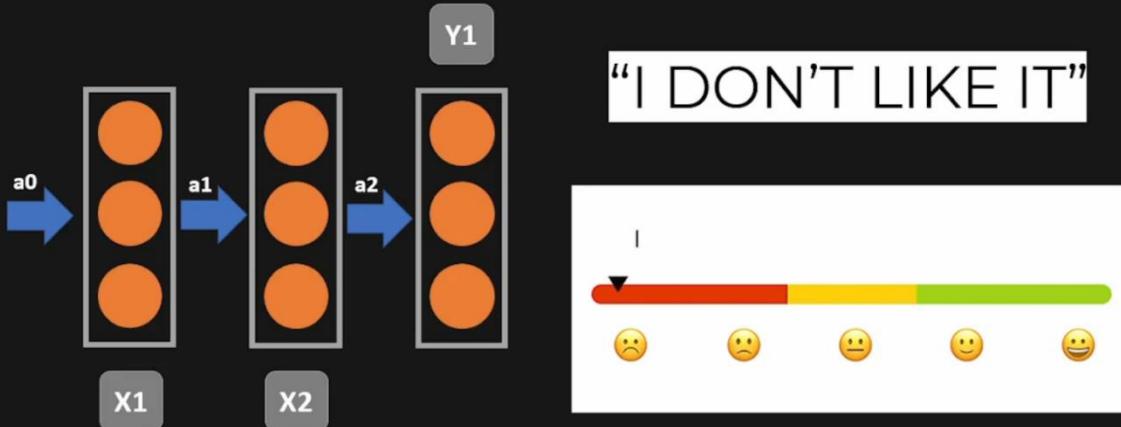
# One to One



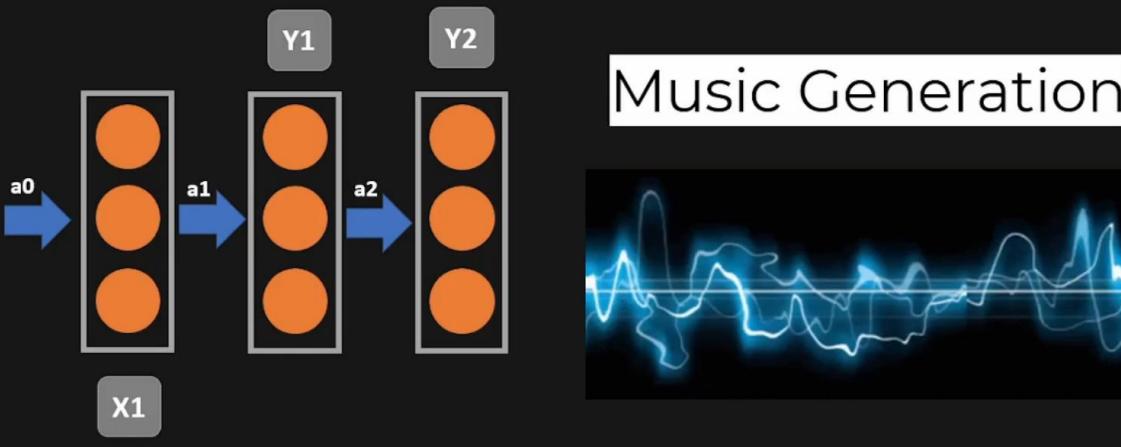
# Many to Many



# Many to One



# One to Many

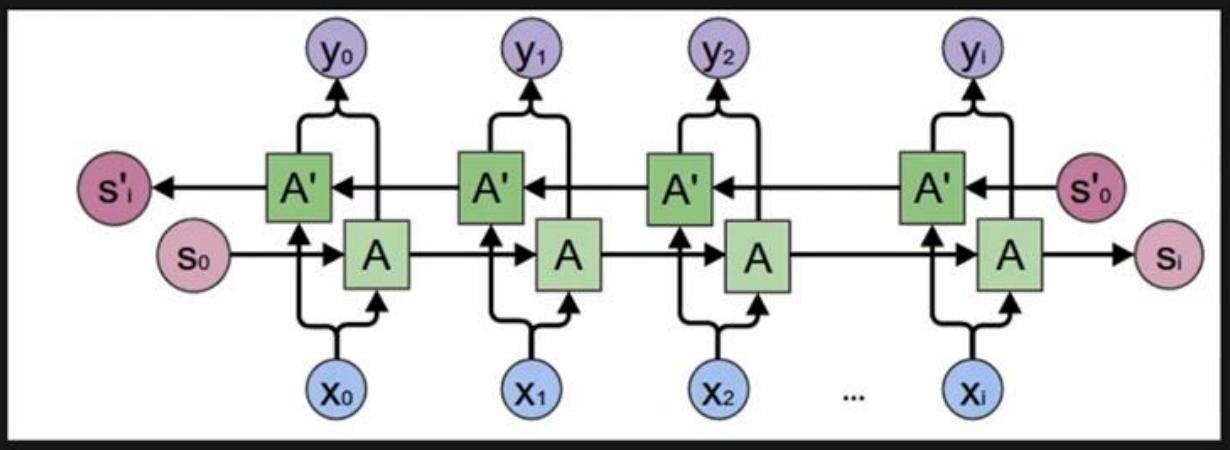
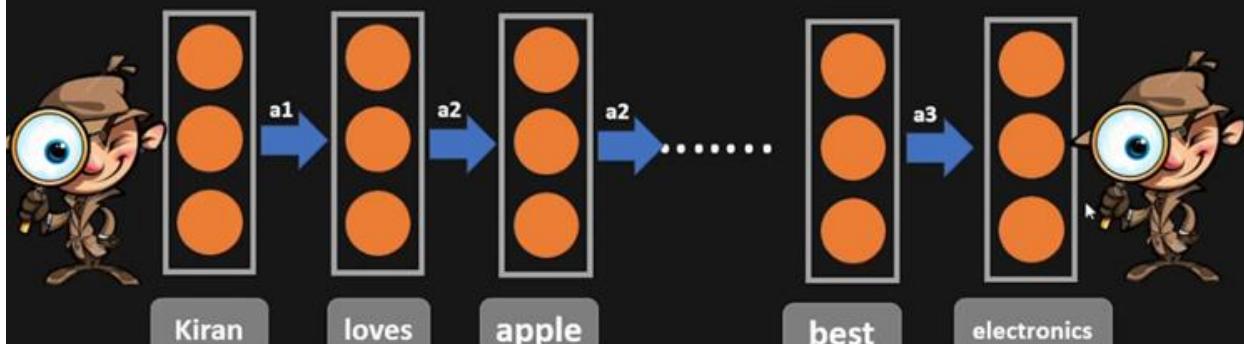


# Bidirectional RNN

## Name Entity Recognition

Kiran loves **apple**, it keeps him **healthy**

Kiran loves **apple**, the company produces best **electronics**



## Long Short-Term Memory (LSTM )

Recurrent Neural Networks work just well with short-term dependencies.

That is when applied to problems like

The colour of parrot is ---

Prediction :

The colour of parrot is Green

Something that was said long before, cannot be recalled when making predictions in the present.

Example

I have 10 years experience in most reputed IT organizations. Last year I switched the company.

.....

I have strong knowledge in ..... technologies.

Prediction :

I have strong knowledge in Microsoft technologies.

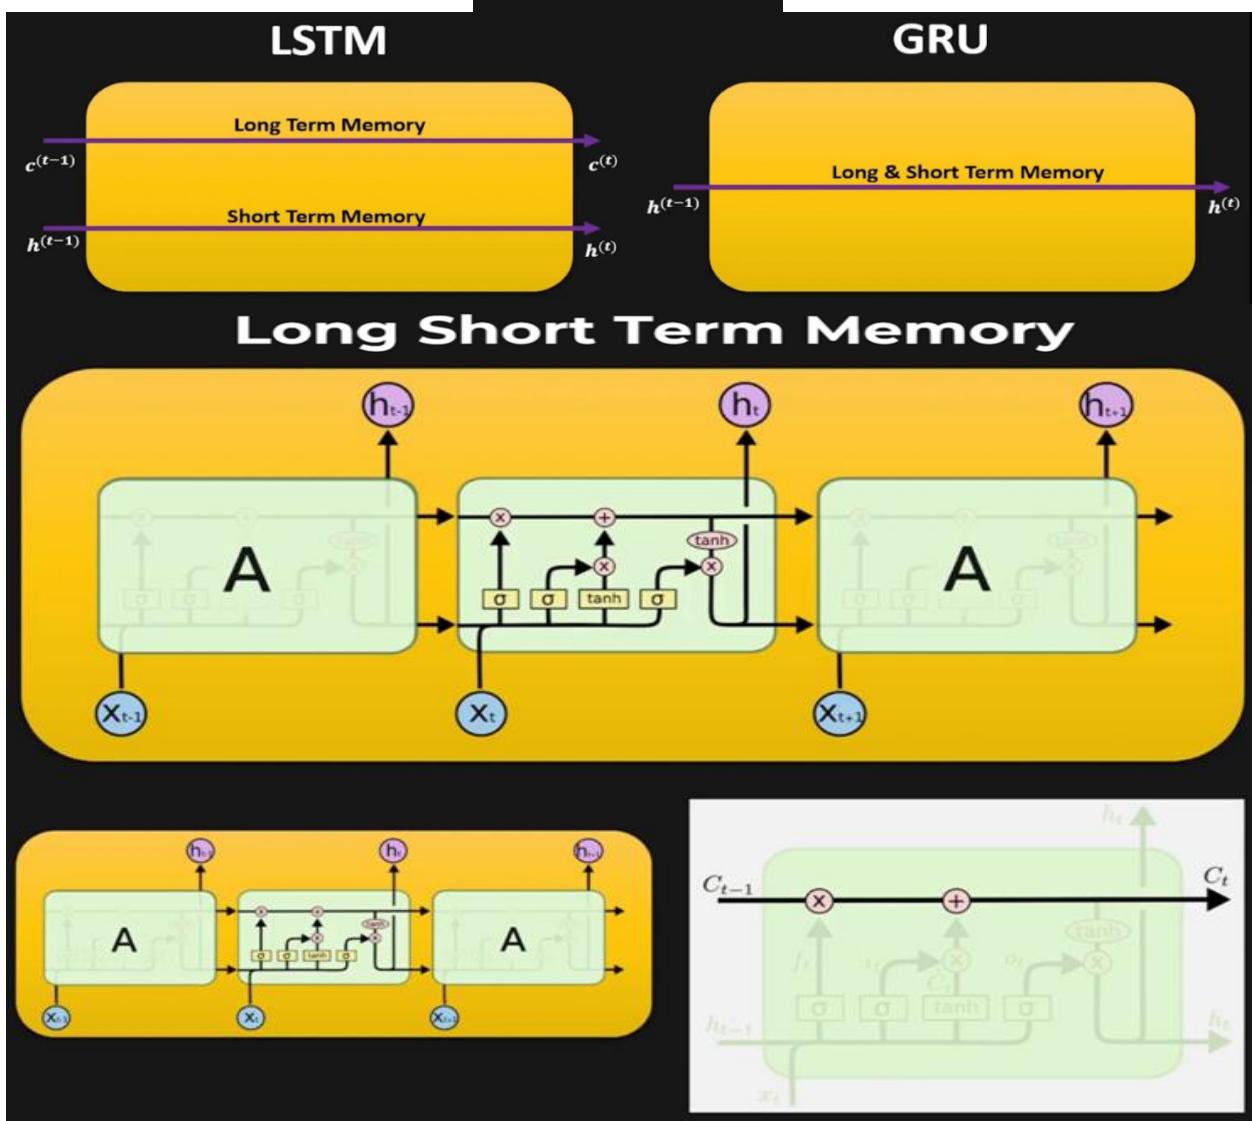
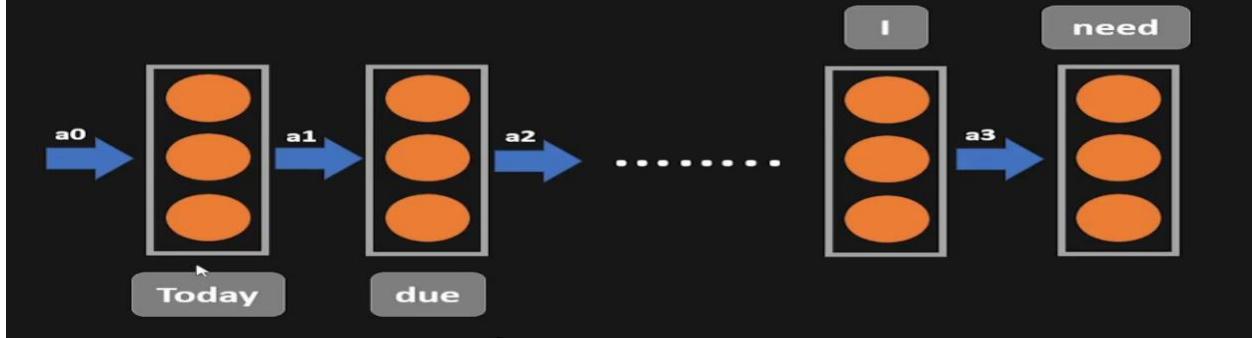
✓RNN suffers from two problems: vanishing gradient and exploding gradient, which make it unusable.

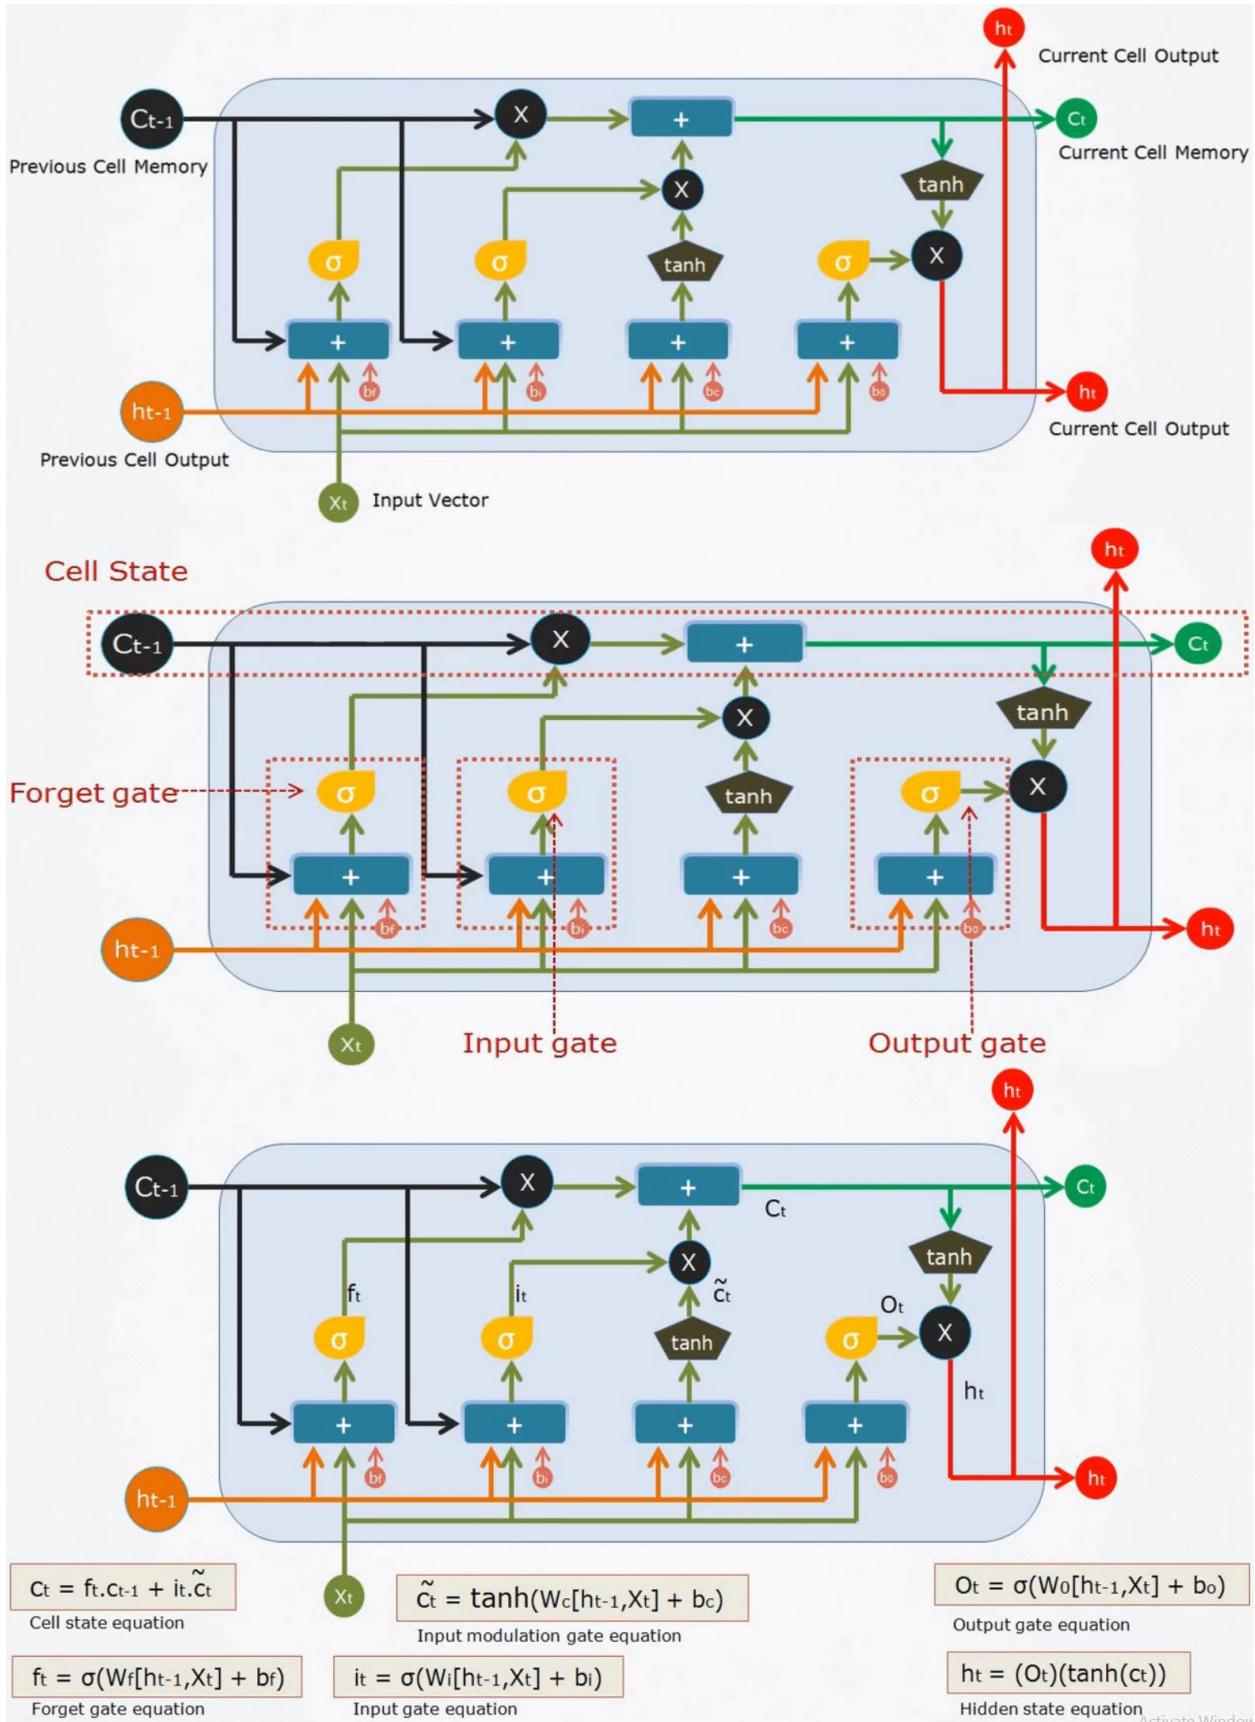
✓LSTM (long short term memory) was invented to solve this issue by explicitly introducing a memory unit, called the cell into the network.

✓This single unit makes decision by considering the current input, previous output and previous memory. And it generates a new output and alters its memory

Today, due to my current job situation and family conditions, I **need** to take a loan.

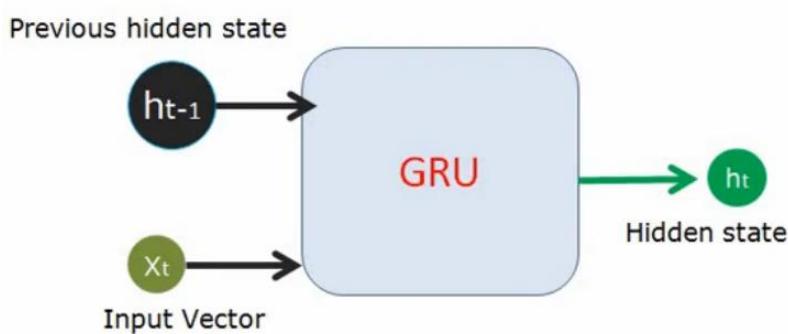
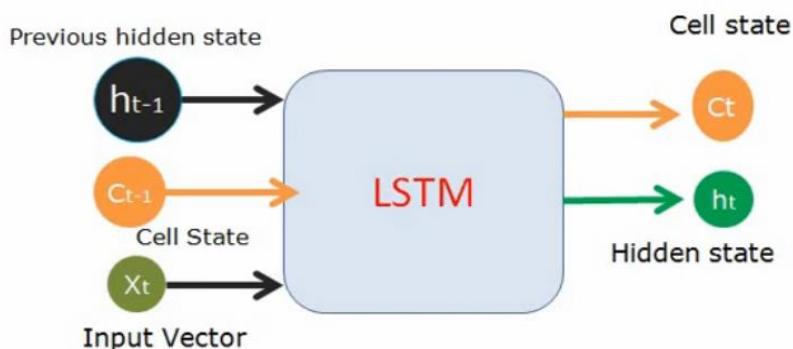
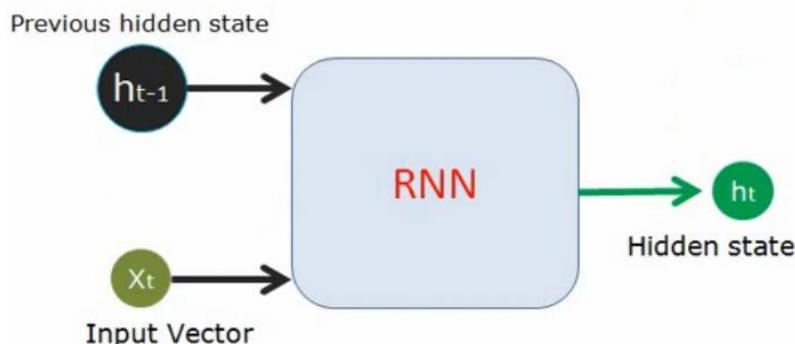
Last Year, due to my current job situation and family conditions, I **had** to take a loan.

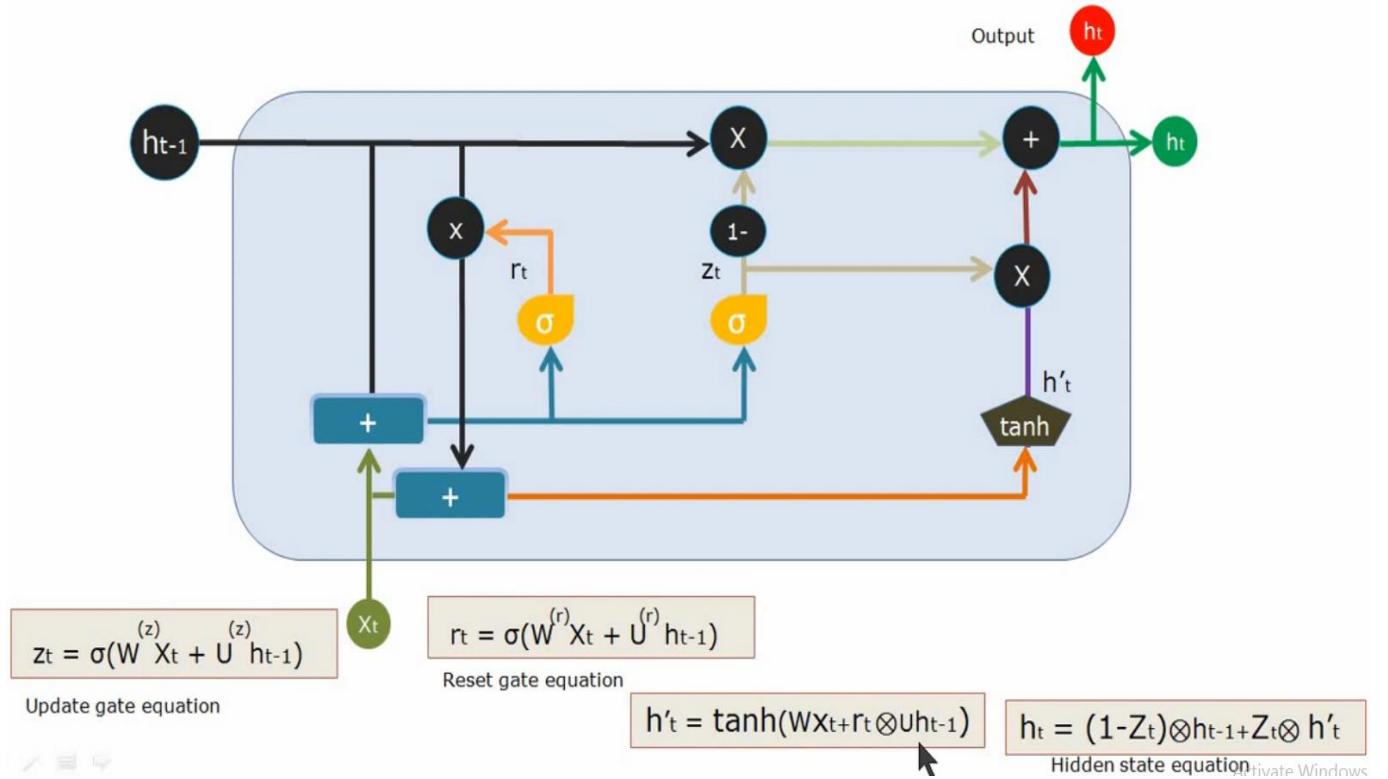
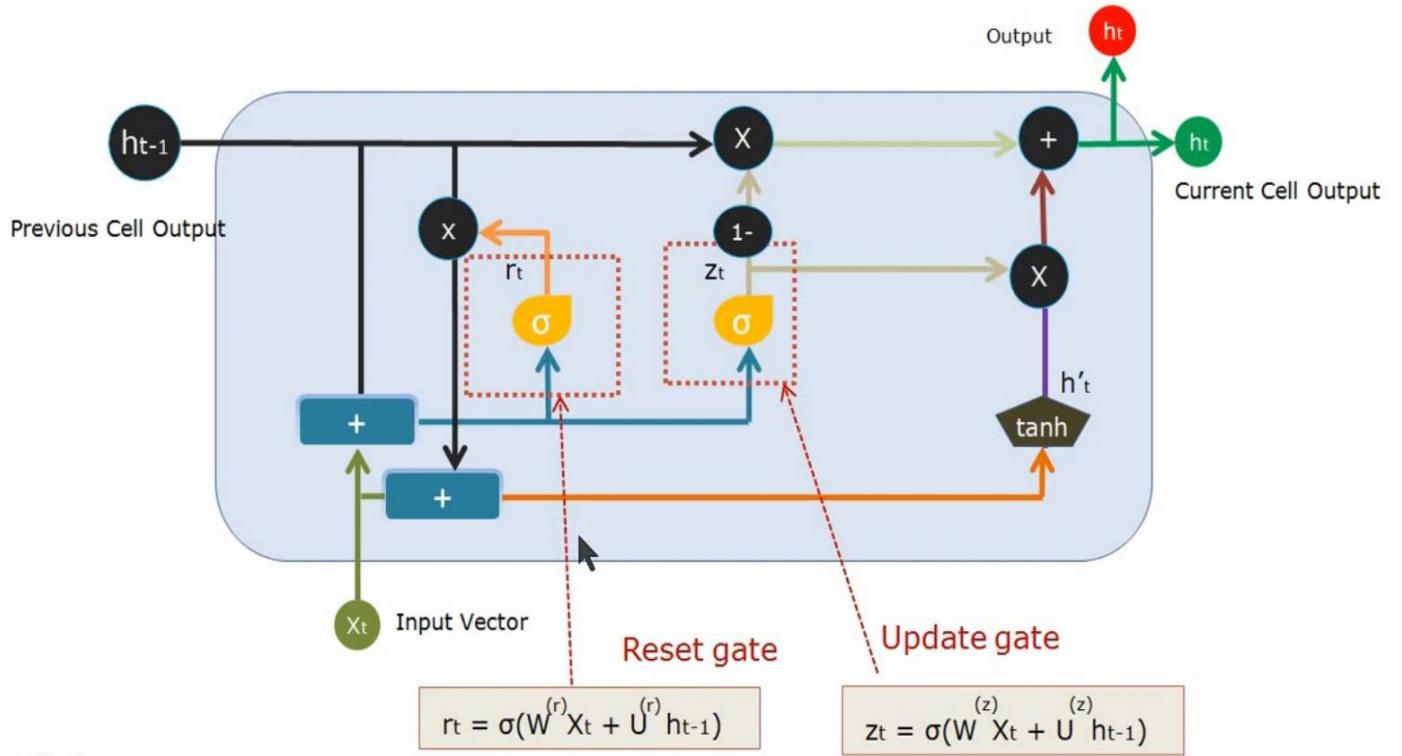




## GRU

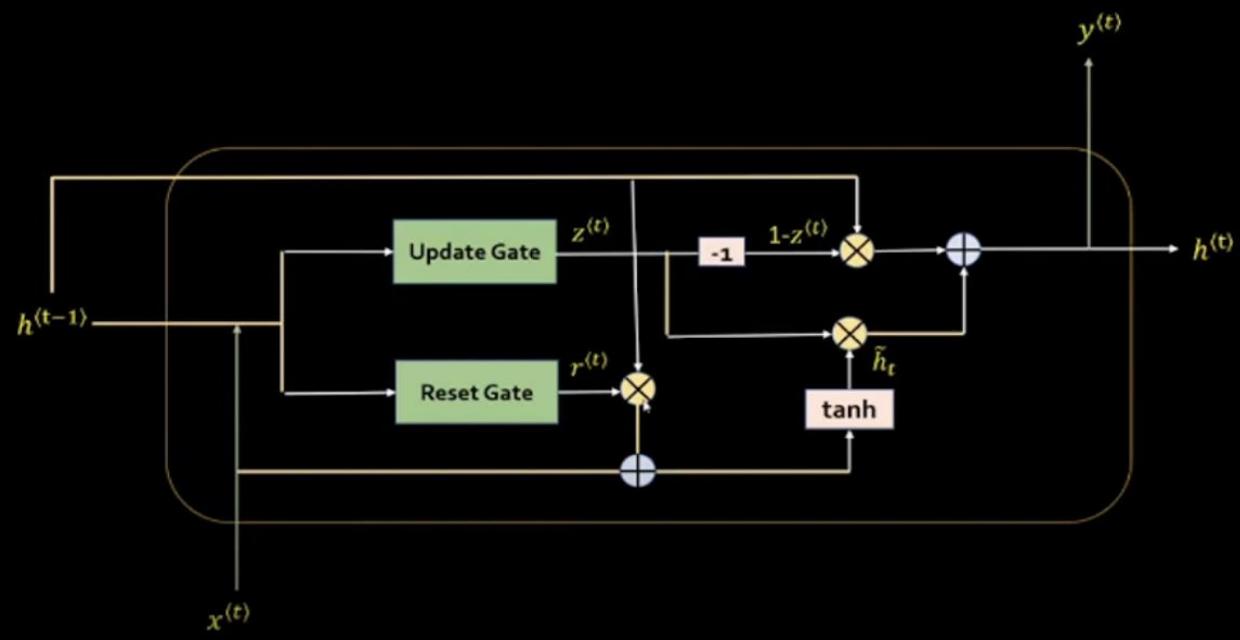
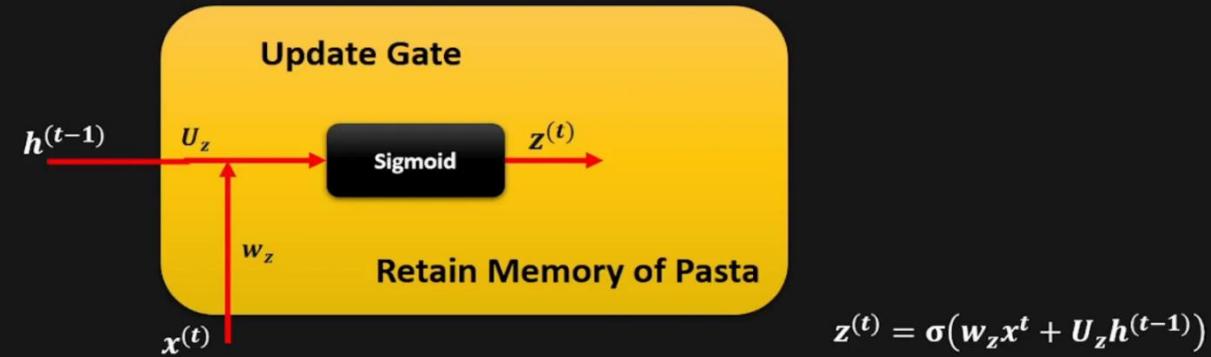
- ✓ Very similar to LSTM
- ✓ It merges the cell state and hidden state.
- ✓ It combines the forget and input gates into a single "update gate".
- ✓ Computationally more efficient.
  - less parameters, less complex structure





# Gated Recurrent Units

Kiran eats **samosa** almost everyday, it shouldn't be hard to guess that his favorite cuisine is **Indian**. His brother Bhavin however is a lover of **pasta** and **cheese** that means Bhavin's favorite cuisine is **Italian**.



```
# First, we get the data
dataset = pd.read_csv('IBM_2006-01-01_to_2018-01-01.csv', index_col='Date', parse_dates=['Date'])
dataset.head()
```

Open High Low Close Volume Name

Date	Open	High	Low	Close	Volume	Name
2006-01-03	82.45	82.55	80.81	82.06	11715200	IBM
2006-01-04	82.20	82.50	81.33	81.95	9840600	IBM
2006-01-05	81.40	82.90	81.00	82.50	7213500	IBM
2006-01-06	83.95	85.03	83.41	84.95	8197400	IBM
2006-01-09	84.10	84.25	83.38	83.73	6858200	IBM

```
[ ] # Checking for missing values
training_set = dataset[:'2016'].iloc[:,1:2].values
test_set = dataset['2017':].iloc[:,1:2].values
```

► We have chosen 'High' attribute for prices. Let's see what it looks like

```
dataset["High"][:'2016'].plot(figsize=(16,4),legend=True)
dataset["High"]['2017':].plot(figsize=(16,4),legend=True)
plt.legend(['Training set (Before 2017)', 'Test set (2017 and beyond)'])
plt.title('IBM stock price')
plt.show()
```



► # Scaling the training set

```
sc = MinMaxScaler(feature_range=(0,1))
training_set_scaled = sc.fit_transform(training_set)
```

► Since GRU store long term memory state, we create a data structure with 60 timesteps and 1 output  
# So for each element of training set, we have 60 previous training set elements

```
X_train = []
y_train = []
for i in range(60,2769):
    X_train.append(training_set_scaled[i-60:i,0])
    y_train.append(training_set_scaled[i,0])
X_train, y_train = np.array(X_train), np.array(y_train)
```

```
[ ] # Reshaping X_train for efficient modelling
X_train = np.reshape(X_train, (X_train.shape[0],X_train.shape[1],1))
```

```

# The GRU architecture
modelGRU = Sequential()
# First GRU layer with Dropout regularisation
modelGRU.add(GRU(units=50, return_sequences=True, input_shape=(X_train.shape[1],1), activation='tanh'))
modelGRU.add(Dropout(0.2))
# Second GRU layer
modelGRU.add(GRU(units=50, return_sequences=True, input_shape=(X_train.shape[1],1), activation='tanh'))
modelGRU.add(Dropout(0.2))
# Third GRU layer
modelGRU.add(GRU(units=50, return_sequences=True, input_shape=(X_train.shape[1],1), activation='tanh'))
modelGRU.add(Dropout(0.2))
# Fourth GRU layer
modelGRU.add(GRU(units=50, activation='tanh'))
modelGRU.add(Dropout(0.2))
# The output layer
modelGRU.add(Dense(units=1))
# Compiling the RNN
modelGRU.compile(optimizer=SGD(lr=0.01, decay=1e-7, momentum=0.9, nesterov=False), loss='mean_squared_error')
# Fitting to the training set
modelGRU.fit(X_train,y_train,epochs=50,batch_size=150)

```

Epoch 1/50  
19/19 [=====] - 10s 223ms/step - loss: 0.1769  
Epoch 2/50  
19/19 [=====] - 4s 227ms/step - loss: 0.0647  
Epoch 3/50  
19/19 [=====] - 4s 230ms/step - loss: 0.0276  
Epoch 48/50  
19/19 [=====] - 5s 239ms/step - loss: 0.0022  
Epoch 49/50  
19/19 [=====] - 5s 239ms/step - loss: 0.0021  
Epoch 50/50  
19/19 [=====] - 5s 244ms/step - loss: 0.0021  
<tensorflow.python.keras.callbacks.History at 0x7f444e914898>

```

# Preparing X_test and predicting the prices
X_test = []
for i in range(60,311):
    X_test.append(inputs[i-60:i,0])
X_test = np.array(X_test)
X_test = np.reshape(X_test, (X_test.shape[0],X_test.shape[1],1))
GRU_predicted_stock_price = modelGRU.predict(X_test)
GRU_predicted_stock_price = sc.inverse_transform(GRU_predicted_stock_price)

# Visualizing the results for LSTM
plot_predictions(test_set,predicted_stock_price)

```



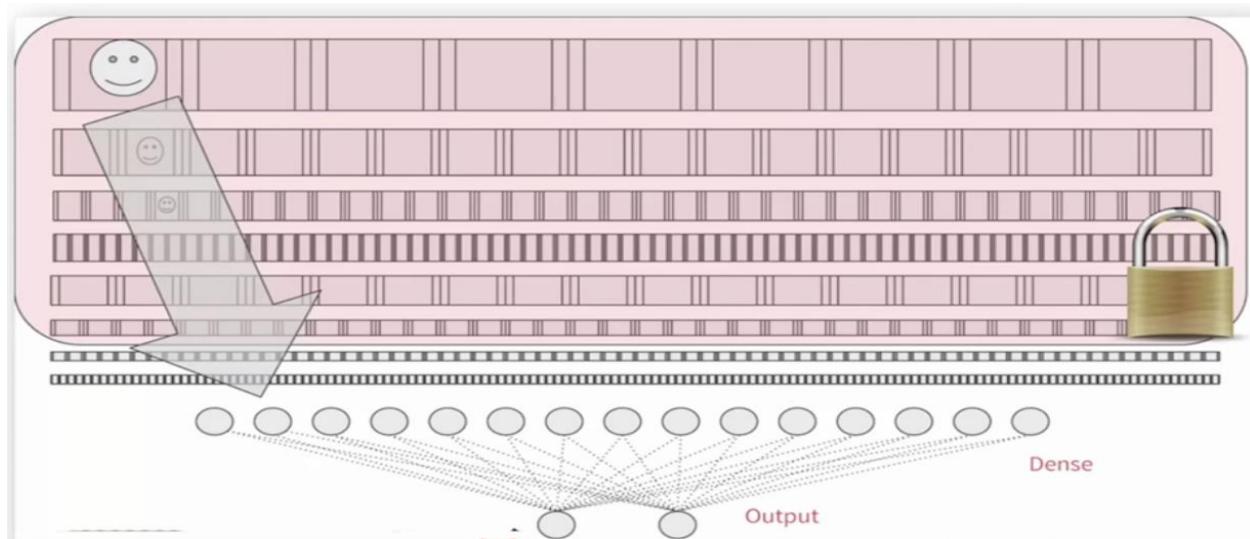
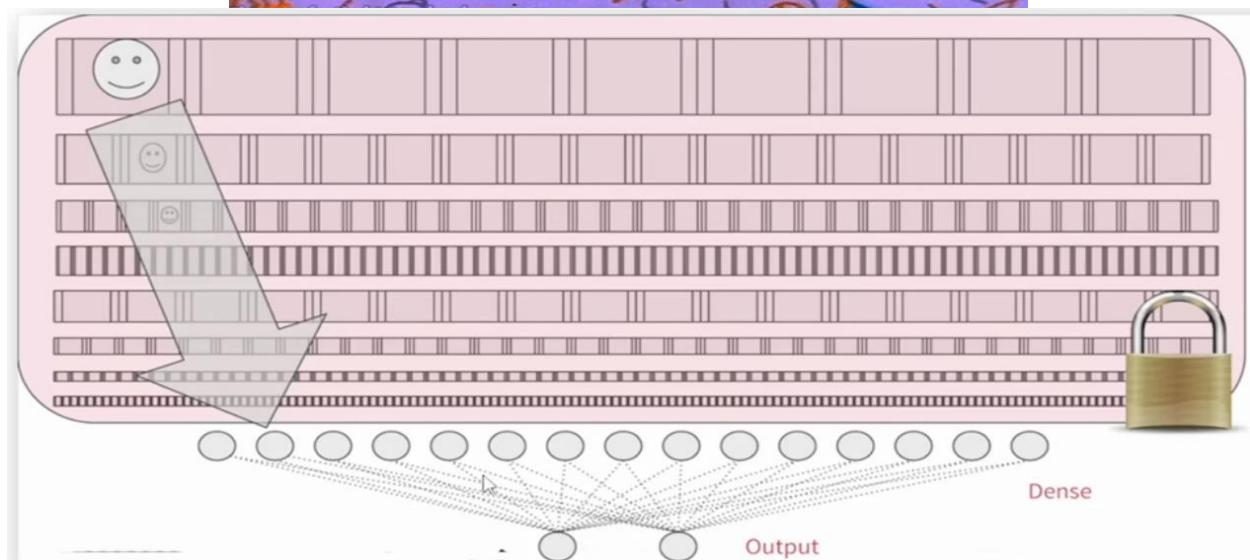
```

# Evaluating our model
return_rmse(test_set,predicted_stock_price)

```

The root mean squared error is 2.592514562101024.

## Transfer Learning



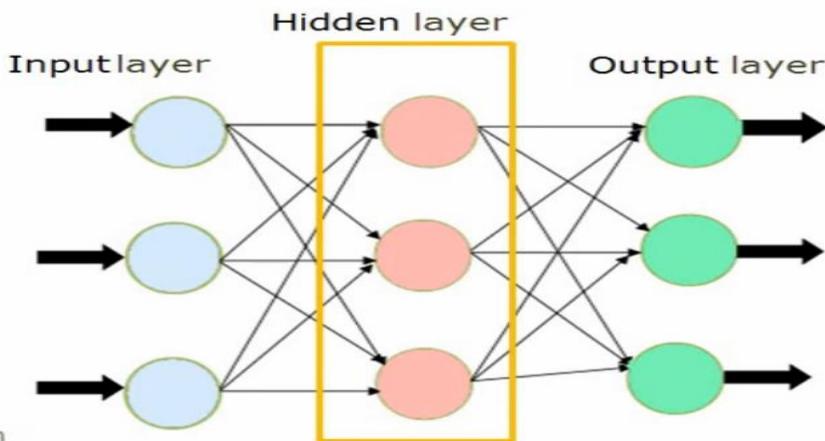
## TensorFlow Version



```
import tensorflow as tf  
  
tf.__version__  
  
'2.3.0'
```

## Keras layers API

Layers are the basic building blocks of neural networks in Keras



```
import tensorflow as tf  
from tensorflow.keras import layers  
  
#creating layer  
layer = layers.Dense(4)  
  
inputs = tf.random.uniform(shape=(5, 10))  
outputs = layer(inputs)  
  
outputs
```



```
<tf.Tensor: shape=(5, 4), dtype=float32, numpy=  
array([[ 0.23317587, -0.5723395 , -0.5922195 , -0.15031755],  
       [ 0.42952067, -1.0333114 ,  0.36681855,  0.4331702 ],  
       [ 0.02772742, -0.38887852, -0.36778417, -0.27738553],  
       [ 0.14641379, -0.73764247,  0.25093132,  0.03845632],  
       [ 0.8514145 , -1.3955642 , -0.01848543,  0.46734422]],  
      dtype=float32)>
```

## Dense Layer

Dense implements the operation:

```
model.add(layers.Dense ())
```

**output = activation(dot(input, kernel) + bias)**

$$Y = \sigma(WX + B)$$

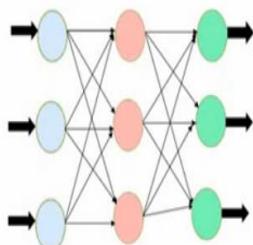
**activation** - activation function

**kernel** is a weights matrix created by the layer

**bias** is a bias vector created by the layer (only applicable if use\_bias True).

## Dense Layer

```
tf.keras.layers.Dense  
(  
    units,  
    activation=None,  
    use_bias=True,  
    kernel_initializer="glorot_uniform",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```



**units**: Positive integer, dimensionality of the output space.  
(Number of Neurons)

**activation**: Activation function to use. If you don't specify anything, no activation is applied  
(ie. "linear" activation:  $a(x) = x$ ).

**use\_bias**: Boolean, whether the layer uses a bias vector.

**kernel\_initializer**: Initializer for the kernel weights matrix.

**bias\_initializer**: Initializer for the bias vector.

**kernel\_regularizer**: Regularizer function applied to the kernel weights matrix.

**bias\_regularizer**: Regularizer function applied to the bias vector.

**activity\_regularizer**: Regularizer function applied to the output of the layer (its "activation").

**kernel\_constraint**: Constraint function applied to the kernel weights matrix.

**bias\_constraint**: Constraint function applied to the bias vector.

## Flatten Layer

Flattens the input. Does not affect the batch size.

```
model.add(layers.Flatten())
```

**Flatten layer** converts the pooled vector to a single column

If inputs are shaped (batch,) without a channel dimension, then flattening adds an extra channel dimension and output shapes are (batch, 1).



## Dropout Layer

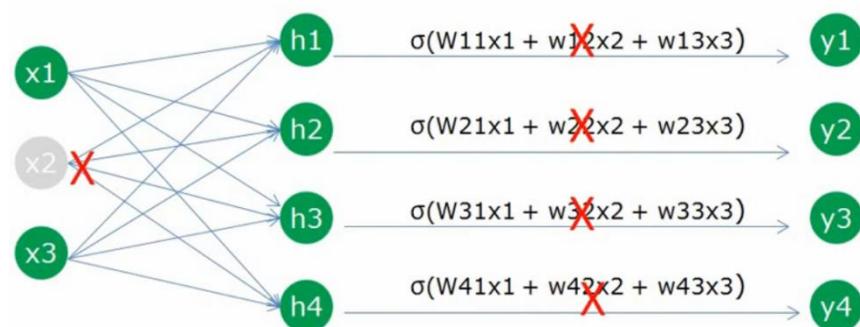
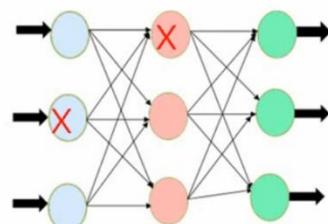
```
model.add(layers.Dropout (rate, noise, _shape=None))
```

**rate** - Float between 0 and 1. Fraction of the input units to drop.

**Dropout** is a technique where randomly selected neurons are ignored during training. They are “dropped-out” randomly.

Use this layer to avoid model overfitting.

Dropout is a technique to regularize in neural networks. When we drop certain nodes out, these units are not considered during a particular forward or backward pass in a network



## Multiple layers

### # Create 3 layers

```
layer1 = layers.Dense(2, activation="relu", name="layer1")
layer2 = layers.Dense(3, activation="relu", name="layer2")
layer3 = layers.Dense(4, name="layer3") # Call layers on a
test input x = tf.ones((3, 3))
y = layer3(layer2(layer1(x)))
```

### # Define Sequential model with 3 layers

```
model = keras.Sequential(
(
    [
        layers.Dense(2, activation="relu", name="layer1"),
        layers.Dense(3, activation="relu", name="layer2"),
        layers.Dense(4, name="layer3"),
    ]
)
x = tf.ones((3, 3))
y = model(x)
```

### # Define Sequential model with 3 layers

```
model = keras.Sequential()
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(3, activation="relu"))
model.add(layers.Dense(4))
x = tf.ones((3, 3))
y = model(x)
```

## Specifying the input shape in advance

### # Define Sequential model with 1 layer

```
model = keras.Sequential()
model.add(keras.Input(shape=(4,)))
model.add(layers.Dense(2, activation="relu"))
model.summary()
```

### # Define Sequential model with 1 layer

```
model = keras.Sequential()
model.add(layers.Dense(2, activation="relu", input_shape=(4,)))
model.summary()
```

Passing an Input object to your model, so that it knows its input shape from the start while creating model.

Note that the Input object is not displayed as part of `model.layers`, since it isn't a layer:

## TensorFlow weights visualization

```
▶ from tensorflow.keras import layers
import tensorflow as tf
model = tf.keras.Sequential()
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(4))
x = tf.ones((3,3))
y = model(x)
model.weights

▷ [<tf.Variable 'dense_30/kernel:0' shape=(3, 2) dtype=float32, numpy=
array([[ 0.11687744, -0.7544714 ],
       [ 0.51128435, -0.55170286],
       [ 0.4792596 , -0.47417873]], dtype=float32)>,
<tf.Variable 'dense_30/bias:0' shape=(2,) dtype=float32, numpy=array([0., 0.], dtype=float32)>,
<tf.Variable 'dense_31/kernel:0' shape=(2, 4) dtype=float32, numpy=
array([[ 0.8916414 ,  0.47133827,  0.43887568,  0.56327534],
       [ 0.89368916, -0.7732425 , -0.7410698 ,  0.32620478]], dtype=float32)>,
<tf.Variable 'dense_31/bias:0' shape=(4,) dtype=float32, numpy=array([0., 0., 0., 0.], dtype=float32)>]
```

## TensorFlow model summary

```
▶ from tensorflow.keras import layers
import tensorflow as tf
model = tf.keras.Sequential()
model.add(layers.Dense(2, activation="relu"))
model.add(layers.Dense(4))
x = tf.ones((3,3))
y = model(x)
model.summary()

▷ Model: "sequential_22"
-----  
Layer (type)          Output Shape         Param #
-----  
dense_49 (Dense)      (3, 2)              8  
-----  
dense_50 (Dense)      (3, 4)              12  
-----  
Total params: 20  
Trainable params: 20  
Non-trainable params: 0
```

The screenshot shows a web browser window displaying the Keras API reference documentation at [keras.io/api/](https://keras.io/api/). The page has a dark-themed header with the Keras logo and navigation links. The main content area is titled "Keras API reference" and contains sections for "Models API" and "Layers API", each with a bulleted list of sub-topics. A search bar is at the top right, and a sidebar on the right lists various API categories. The bottom of the screen shows the Windows taskbar with several pinned icons and the system tray.

Keras API reference

Models API

- The Model class
- The Sequential class
- Model training APIs
- Model saving & serialization APIs

Layers API

- The base Layer class
- Layer activations
- Layer weight initializers
- Layer weight regularizers
- Layer weight constraints
- Core layers
- Convolution layers
- Pooling layers
- Recurrent layers
- Preprocessing layers
- Normalization layers
- Regularization layers
- Attention layers
- Padding layers

Code examples

Search Keras documentation...

About Keras

Getting started

Developer guides

Keras API reference

Models API

Layers API

Callbacks API

Data preprocessing

Optimizers

Metrics

Losses

Built-in small datasets

Keras Applications

Utilities

Keras API reference

Models API

Layers API

Callbacks API

Data preprocessing

Optimizers

Metrics

Losses

Built-in small datasets

Keras Applications

Utilities

keras.io/api/

Windows Taskbar icons: File Explorer, Edge, File Manager, Task View, File History, OneDrive, Google Chrome, File Explorer, Notepad, Coaching.com

System Tray: Battery, Signal, Volume, Date: 17-03-2020, Time: 07:58

## Why do we need to save model?

```
from pandas import read_csv
# load the dataset
df = read_csv('https://rawexample.com/output.csv', header=None)

# split into input and output columns
X, y = df.values[:, :-1], df.values[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)

# determine the number of input features
n_features = X_train.shape[1]

# define model
model = Sequential()
model.add(Dense(10, activation='relu', kernel_initializer='he_normal',
input_shape=(n_features,)))
model.add(Dense(1, activation='sigmoid'))

# compile the model
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# fit the model
model.fit(X_train, y_train, epochs=150, batch_size=32, verbose=0)

# make a prediction
PredictedValue = model.predict([1,0,0.99539,0.18641,-0.45300])
print('Predicted: %.3f' % PredictedValue)
```

## Scikit-Learn Machine Learning Models

### Save The Model

```
from sklearn.linear_model import LogisticRegression
import pickle
model = LogisticRegression()
model.fit(xtrain, ytrain)

# save the model to disk
pickle.dump(model, open("modelName.pkl", 'wb'))
```

### Load The Model

```
model = pickle.load(open("modelName.pkl", 'rb'))
Result = model.predict(InputValue)
```

### Use **Pickle** to serialise and save the models

### Save The Model

```
from sklearn.linear_model import LogisticRegression
from sklearn.externals import joblib
model = LogisticRegression()
model.fit(xtrain, ytrain)

# save the model to disk
joblib.dump(model, "modelName.pkl")
```

### Load The Model

```
model = joblib.load ("modelName.pkl")
Result = model.predict(InputValue)
```

### Use **joblib** to serialise and save the models

### Saving model architecture by two ways

Save Model to JSON.

Save Model to YAML.

### Save The Model

```
from keras.models import Sequential
from keras.layers import Dense# create model

model = Sequential()
# Fit the model
model.fit(xtrain, ytrain)

# serialize to JSON
json_file = model.to_json()
with open ("modelName.json", "w") as file:
file.write(json_file)
# serialize weights to HDF5
model.save_weights("modelName.h5")
```

### Load The Model

```
from keras.models import model_from_json

# load json and create model
file = open(("modelName.json", 'r'))
loaded_model_json = file.read()
file.close()

model = model_from_json(loaded_model_json)

# load weights
model.load_weights ("modelName.h5")

Result = model.predict(InputValue)
```

## Save The Model

```
from keras.models import Sequential  
from keras.layers import Dense# create model  
  
model = Sequential()  
# Fit the model  
model.fit(xtrain, ytrain)  
  
# serialize to JSON  
json_file = model.to_yaml()  
with open("modelName.yaml", "w") as file:  
file.write(json_file)  
  
# serialize weights to HDF5  
model.save_weights("modelName.h5")
```

## Save The Model

```
from keras.models import Sequential  
from keras.layers import Dense# create model  
  
model = Sequential()  
  
# Fit the model  
model.fit(xtrain, ytrain)  
  
# serialize weights to HDF5  
model.save ("modelName.h5")
```

## Load The Model

```
from keras.models import model_from_yaml  
  
# load yaml and create model  
file = open("modelName.yaml", 'r')  
loaded_model_yaml = file.read()  
file.close()  
  
model = model_from_yaml(loaded_model_yaml)  
  
# load weights  
model.load_weights("modelName.h5")  
  
Result = model.predict(InputValue)
```

## Load The Model

```
from keras.models import load_model  
  
# load model  
model = load_model("modelName.h5")  
  
Result = model.predict(InputValue)
```